

# Introdução às Redes de Comunicação

## 2019/2020

### T03

### Application layer

Jorge Granjal  
University of Coimbra



# T03

## Application Layer

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

# T03

# Application Layer

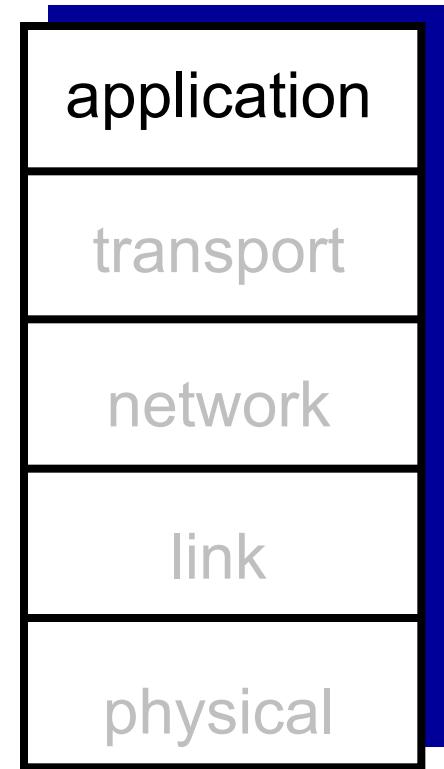
---

## our goals:

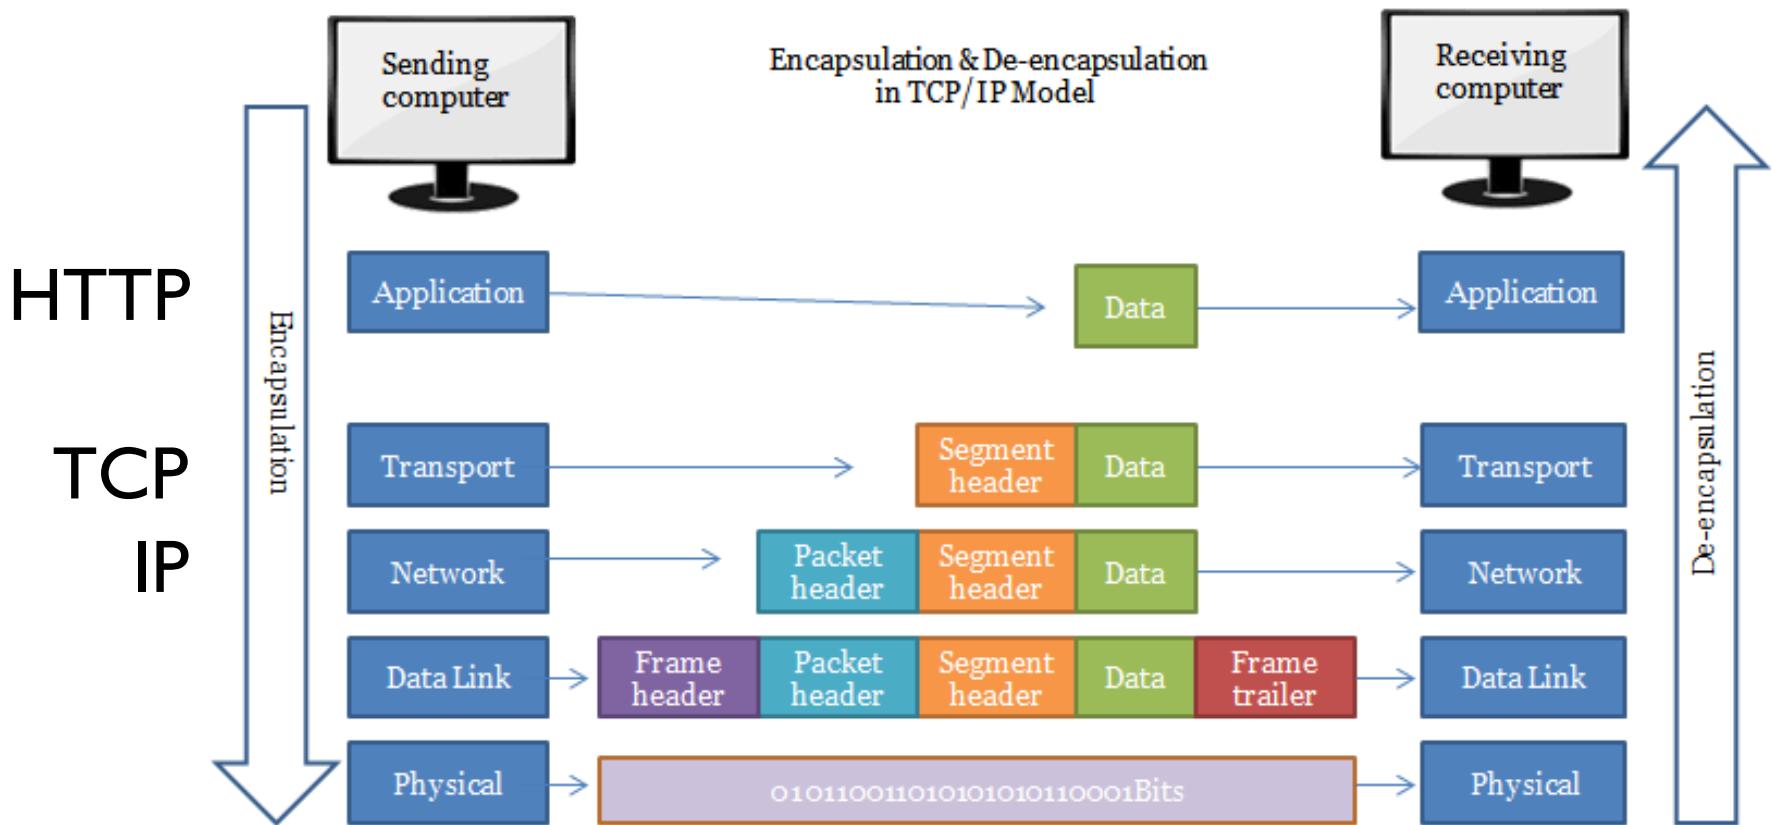
- conceptual, implementation aspects of network application protocols
  - transport-layer service models
  - client-server paradigm
  - peer-to-peer paradigm
  - content distribution networks
- learn about protocols by examining popular application-level protocols
  - HTTP
  - FTP
  - SMTP / POP3 / IMAP
  - DNS

# Internet (TCP/IP) protocol stack

- *application*: supporting network applications
  - FTP, SMTP, HTTP, ...
- *transport*: process-process data transfer
  - TCP, UDP
- *network*: routing of datagrams from source to destination
  - IP, routing protocols
- *link*: data transfer between neighboring network elements
  - Ethernet, 802.111 (WiFi), PPP
- *physical*: bits “on the wire”



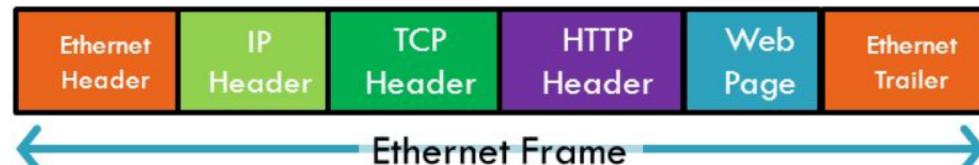
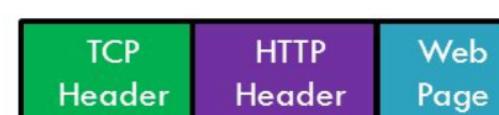
# Encapsulation (yes, again!)



# Encapsulation (yes, again!)

0

| 32 Bit                  |               |                 |              |
|-------------------------|---------------|-----------------|--------------|
| Version                 | Header Length | Type of Service | Total Length |
| Fragment Identification | Flags         | Fragment Offset |              |
| TTL                     | Protocol      | Header Checksum |              |
| Source Address          |               |                 |              |
| Destination Address     |               |                 |              |
| Options & Padding       |               |                 |              |



# Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

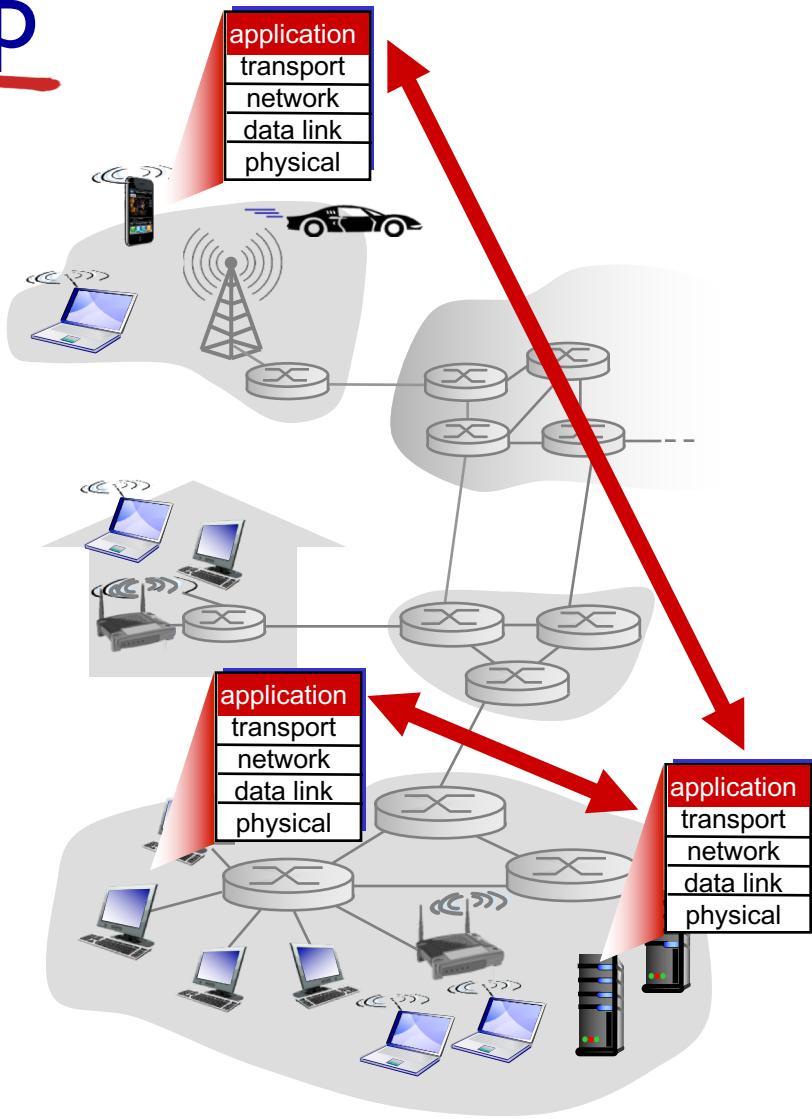
# Creating a network app

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications (operate at the network-layer and below)
- applications on end systems allows for rapid app development, propagation

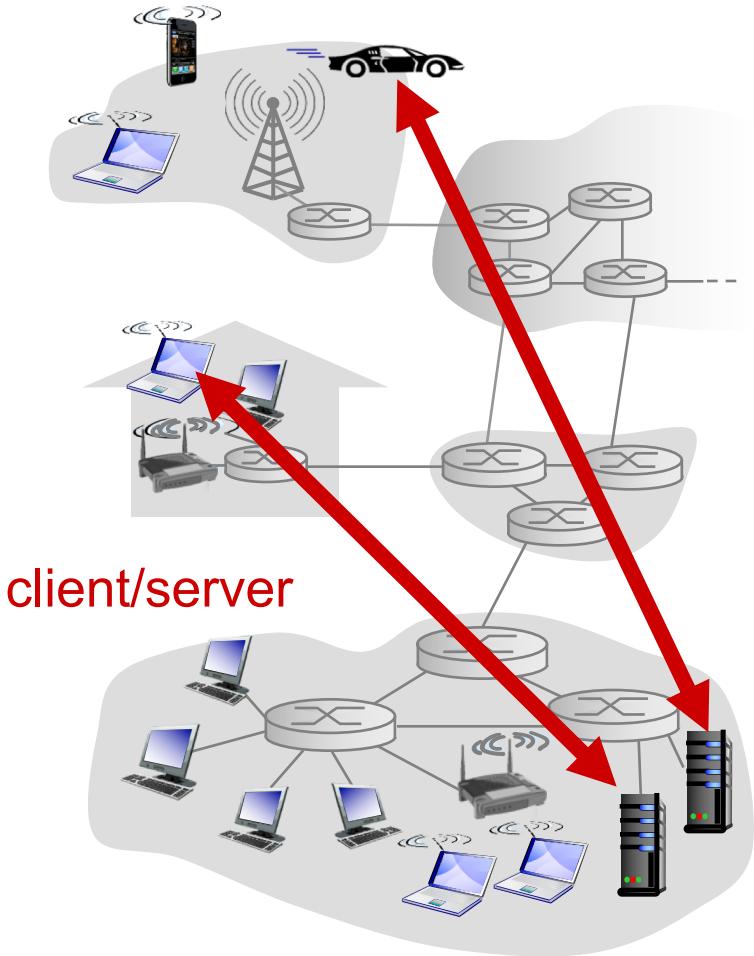


# Application architectures

possible structure of applications:

- client-server
- peer-to-peer (P2P)

# Client-server architecture



## server:

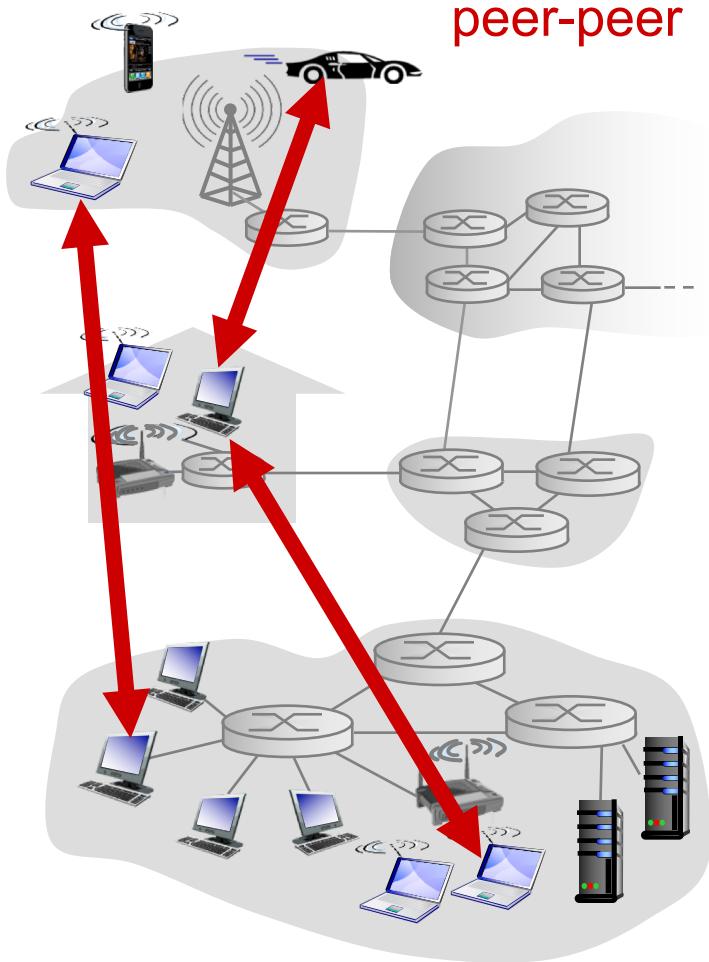
- always-on host
- permanent IP address
- data centers for scaling

## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other
- Examples: Web, FTP, SSH, e-mail

# P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management
- Example: BitTorrent, Skype, IPTV,  
...
- Also, some applications use *hybrid architectures* (e.g. in messaging)



# Processes communicating

*process*: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**.
- IP communications may also take place within same host (via loopback interface, 127.0.0.1)

clients, servers

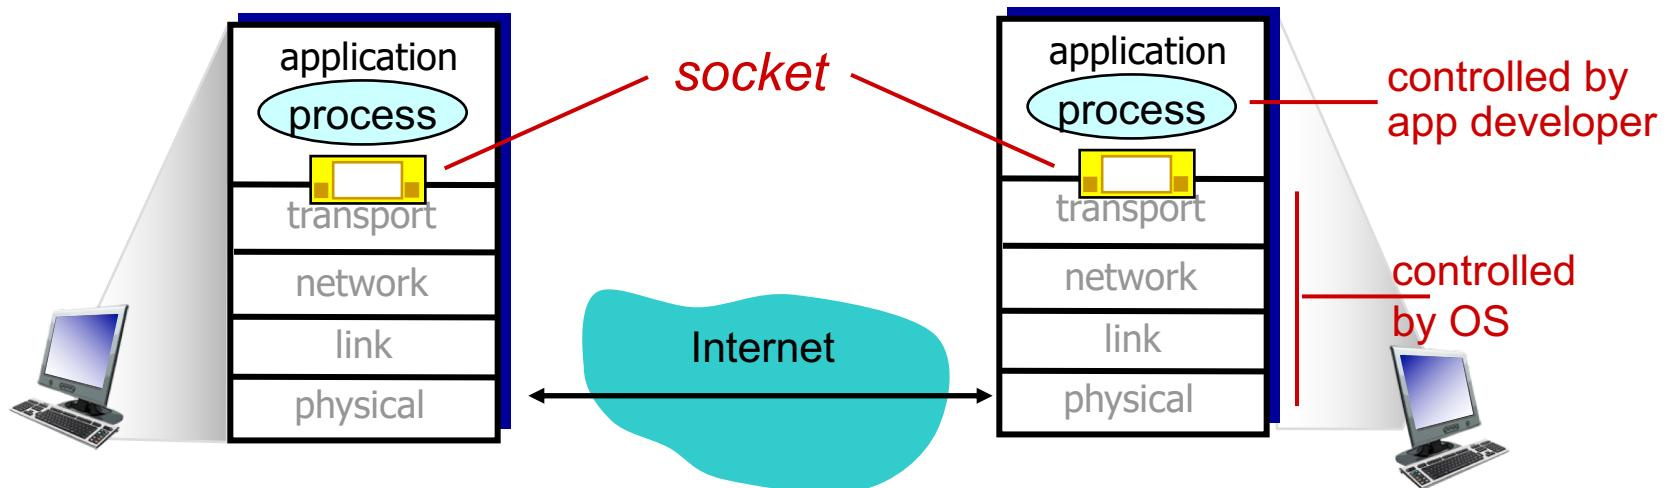
*client process*: process that initiates communication

*server process*: process that waits to be contacted

- aside: applications with P2P architectures have client processes & server processes
- In some applications a process can be both a client and a server

# Sockets

- process sends/receives messages to/from its **socket**
- A socket is the software **interface between the process and the computer network** (between the application layer and the transport layer)
- Is also referred to as the **Application Programming Interface (API)** between the application and the network
- The application chooses the transport protocol (e.g. UDP or TCP) to use the transport-layer services provided by the protocol



# Addressing processes

- to receive messages, processes must have *identifier*
- host device has unique 32-bit IP address
- *Q:* does IP address of host on which process runs suffice for identifying the process?
  - *A:* no, many processes can be running on same host
- *identifier* includes both IP address and port numbers associated with process (the socket) on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - IP address: 128.119.245.12
  - port number: 80
- more shortly...

# Addressing processes

- Popular applications use assigned port numbers, a few examples:

| Port #  | Application Layer Protocol | Type    | Description                                   |
|---------|----------------------------|---------|---|
| 20      | FTP                        | TCP     | File Transfer Protocol - data                 |
| 21      | FTP                        | TCP     | File Transfer Protocol - control              |
| 22      | SSH                        | TCP/UDP | Secure Shell for secure login                 |
| 23      | Telnet                     | TCP     | Unencrypted login                             |
| 25      | SMTP                       | TCP     | Simple Mail Transfer Protocol                 |
| 53      | DNS                        | TCP/UDP | Domain Name Server                            |
| 67/68   | DHCP                       | UDP     | Dynamic Host                                  |
| 80      | HTTP                       | TCP     | HyperText Transfer Protocol                   |
| 123     | NTP                        | UDP     | Network Time Protocol                         |
| 161,162 | SNMP                       | TCP/UDP | Simple Network Management Protocol            |
| 389     | LDAP                       | TCP/UDP | Lightweight Directory Authentication Protocol |
| 443     | HTTPS                      | TCP/UDP | HTTP with Secure Socket Layer                 |

- Full port number and service name list available in the file `/etc/services` in Linux and at <http://www.iana.org>

# App-layer protocol defines

- types of messages exchanged,
  - e.g., request, response
- message syntax:
  - what fields in messages & how fields are delineated
- message semantics
  - meaning of information in fields
- rules for when and how processes send & respond to messages

**open protocols:**

- defined in RFCs
- allows for interoperability
- e.g., HTTP (HyperText Transfer Protocol), SMTP (Simple Mail Transfer Protocol)

**proprietary protocols:**

- e.g., Skype

# What transport service does an app need?

## data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective” (e.g. multimedia)
- other apps (“elastic apps”) make use of whatever throughput they get (e.g. e-mail, web, file transfer)

## security

- encryption, data integrity, ...

*“One can never be too thin or too rich...or have too much throughput!”*

# Transport service requirements: common apps

| <b>application</b>    | <b>data loss</b> | <b>throughput</b>                        | <b>time sensitive</b>          |
|-----------------------|------------------|--|--------------------------------|
| file transfer         | no loss          | elastic                                  | no                             |
| e-mail                | no loss          | elastic                                  | no                             |
| Web documents         | no loss          | elastic                                  | no                             |
| real-time audio/video | loss-tolerant    | audio: 5kbps-1Mbps<br>video:10kbps-5Mbps | yes, 100' s msec               |
| stored audio/video    | loss-tolerant    | same as above                            |                                |
| interactive games     | loss-tolerant    | few kbps up                              | yes, few secs                  |
| text messaging        | no loss          | elastic                                  | yes, 100' s msec<br>yes and no |

# Internet transport protocols services

## TCP service:

- ***reliable transport*** between sending and receiving process (data delivered without errors and in the correct order)
- ***flow control***: sender won't overwhelm receiver
- ***congestion control***: throttle sender when network overloaded
- ***does not provide***: timing, minimum throughput guarantee, security
- ***connection-oriented***: setup required between client and server processes

## UDP service:

- A “*no-frills*” *lightweight transport protocol with minimal services*
- ***unreliable data transfer*** between sending and receiving process
- ***does not provide***: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup
- A programmer need to decide to use UDP or TCP...  
**Q**: why bother? Why is there a UDP?

# Internet apps: application, transport protocols

|                        | <b>application</b> | <b>application<br/>layer protocol</b>   | <b>underlying<br/>transport protocol</b> |
|------------------------|--------------------|---|--|
|                        | e-mail             | SMTP [RFC 2821]                         | TCP                                      |
| remote terminal access |                    | Telnet [RFC 854]                        | TCP                                      |
|                        | Web                | HTTP [RFC 2616]                         | TCP                                      |
|                        | file transfer      | FTP [RFC 959]                           | TCP                                      |
| streaming multimedia   |                    | HTTP (e.g., YouTube),<br>RTP [RFC 1889] | TCP or UDP                               |
| Internet telephony     |                    | SIP, RTP, proprietary<br>(e.g., Skype)  | TCP or UDP                               |

# Securing TCP

## TCP & UDP

- no encryption
- cleartext passwords sent into socket traverse Internet in cleartext

## SSL

- provides encrypted TCP connection
- data integrity
- end-point authentication

## SSL is at app layer

- apps use SSL libraries, that “talk” to TCP

## SSL socket API

- cleartext passwords sent into socket traverse Internet encrypted
- more about network security later in the course!..



# T03

## Application Layer

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

# Web and HTTP

*First, a review...*

- *web page* consists of *objects*
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

www.someschool.edu/someDept/pic.gif

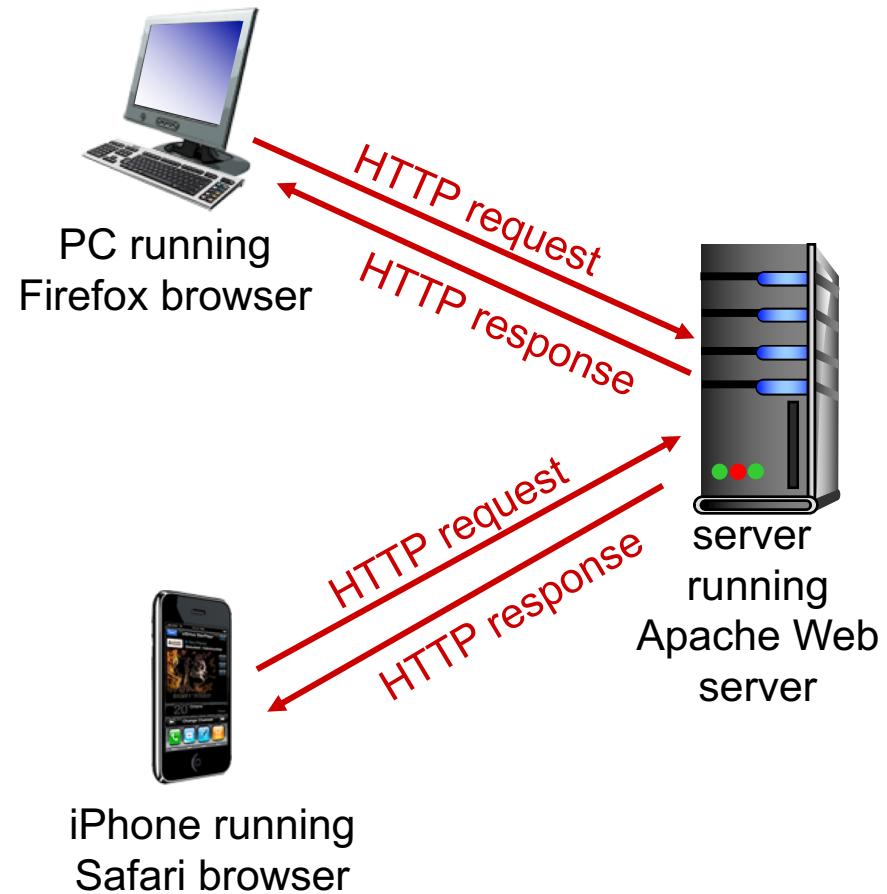
host name

path name

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - **client:** browser that requests, receives, (using HTTP protocol) and "displays" Web objects
  - **server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview (continued)

*uses TCP:*

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

*HTTP is “stateless”*

- server maintains **no** information about past client requests

*aside*

protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP connections

## *non-persistent HTTP*

- at most one object sent over TCP connection
  - connection then closed
- downloading multiple objects required multiple connections

## *persistent HTTP*

- multiple objects can be sent over single TCP connection between client, server
- (is the default mode)

# Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

Ia. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

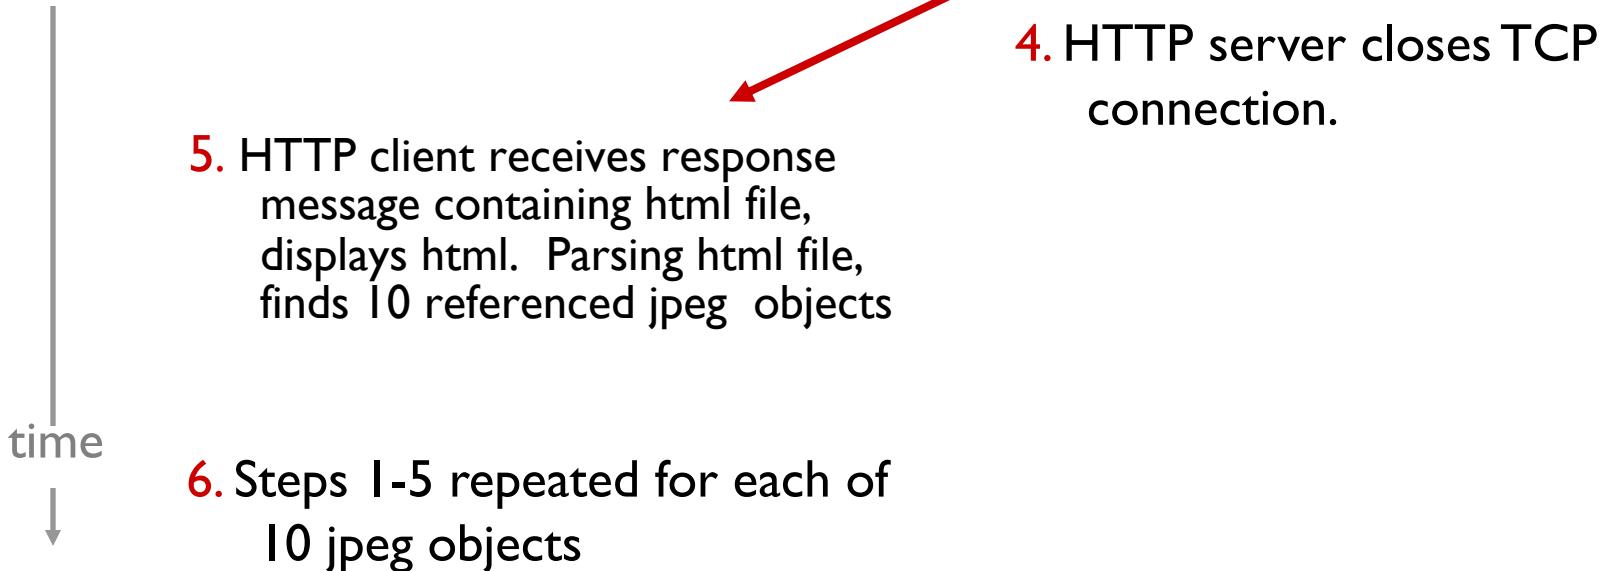
Ib. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time ↓

# Non-persistent HTTP (cont.)

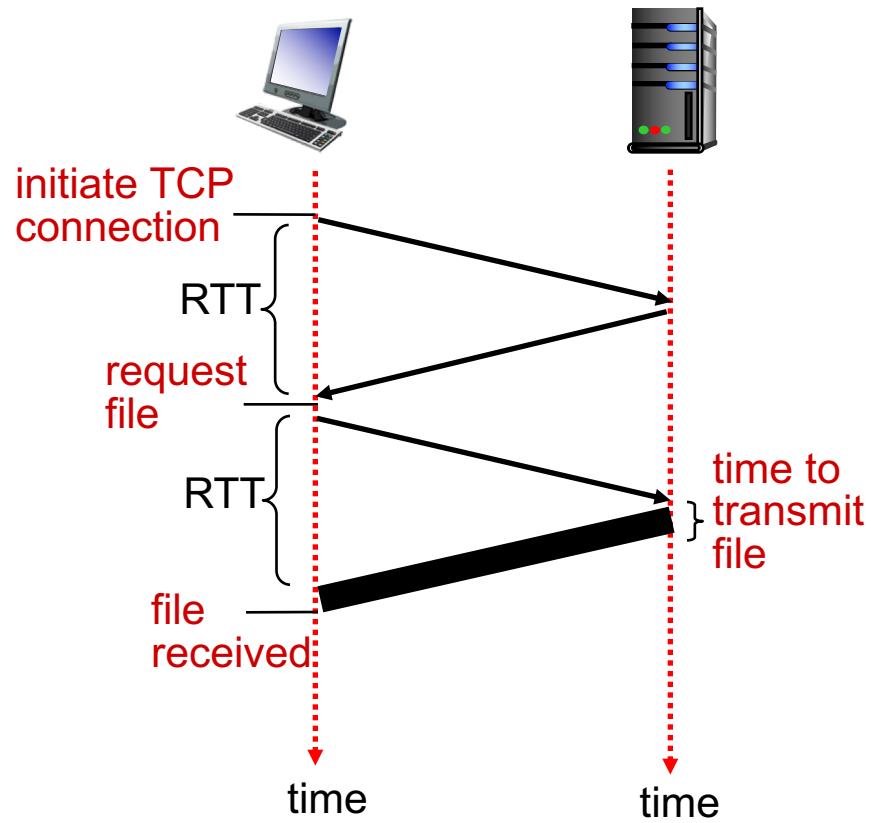


# Non-persistent HTTP: response time

**RTT (definition):** time for a small packet to travel from client to server and back (considering also delays)

**HTTP response time:**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =  
$$2\text{RTT} + \text{file transmission time}$$



# Persistent HTTP

## *non-persistent HTTP issues:*

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## *persistent HTTP:*

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects
- An entire web page (or various) may be sent over a *single (persistent)* TCP connection

# HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:** written in ordinary ASCII (human-readable format) text

request line  
(GET, POST,  
HEAD commands)

header  
lines

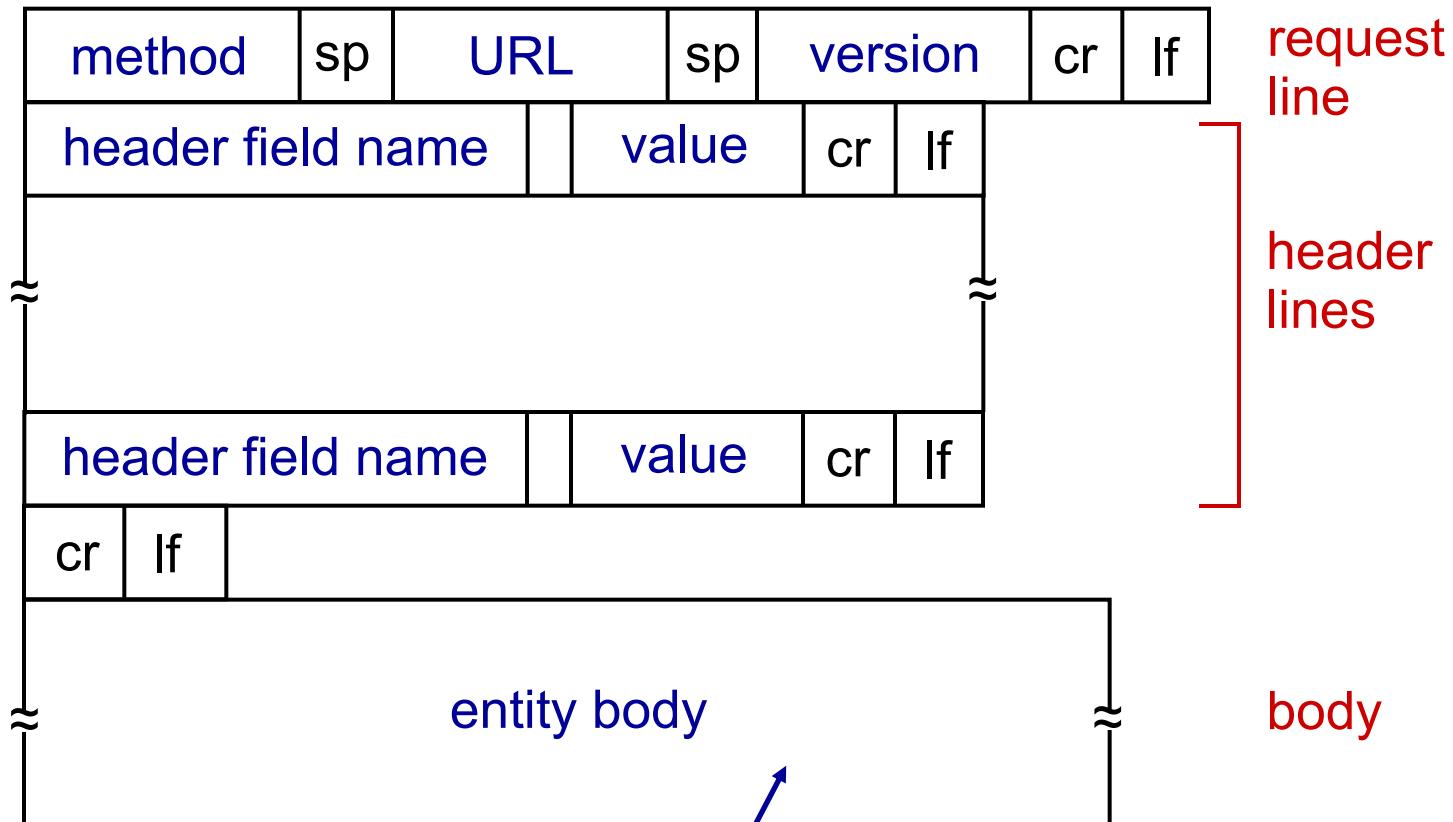
carriage return,  
line feed at start  
of line indicates  
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character  
line-feed character

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

# HTTP request message: general format



Empty with GET, used with POST (e.g. contents  
of a form filled by user)

# Uploading form input

## POST method:

- web page often includes form input
- input is uploaded to server in entity body

## URL method:

- uses GET method
- input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

# Method types

## HTTP/1.0:

- GET
- POST
- HEAD
  - asks server to leave requested object out of response (frequently used for debugging)

## HTTP/1.1:

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

# HTTP response message

status line

(protocol

status code

status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\nDate: Sun, 26 Sep 2010 20:09:20 GMT\r\nServer: Apache/2.0.52 (CentOS)\r\nLast-Modified: Tue, 30 Oct 2007 17:00:02  
GMT\r\nETag: "17dc6-a5c-bf716880"\r\nAccept-Ranges: bytes\r\nContent-Length: 2652\r\nKeep-Alive: timeout=10, max=100\r\nConnection: Keep-Alive\r\nContent-Type: text/html; charset=ISO-8859-  
1\r\n\r\ndata data data data data ...
```

Confirms server is using persistent TCP connection

\* Check out the online interactive exercises for more  
examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

# HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

## **200 OK**

- request succeeded, requested object later in this msg

## **301 Moved Permanently**

- requested object moved, new location specified later in this msg (using the Location: header)

## **400 Bad Request**

- request msg not understood by server

## **404 Not Found**

- requested document not found on this server

## **505 HTTP Version Not Supported**

# Trying out HTTP (client side) for yourself

## I. Telnet to your favorite Web server:

**telnet www.uc.pt 80**

{ opens TCP connection to port 80  
(default HTTP server port)  
at www.uc.pt.  
anything typed in will be sent  
to port 80 at www.uc.pt

## 2. type in a GET HTTP request:

**GET /index.html HTTP/1.1**

**Host: www.uc.pt**

{ by typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

## 3. look at response message sent by HTTP server! (or use Wireshark to look at captured HTTP request/response)

# User-server state: cookies

many Web sites use cookies

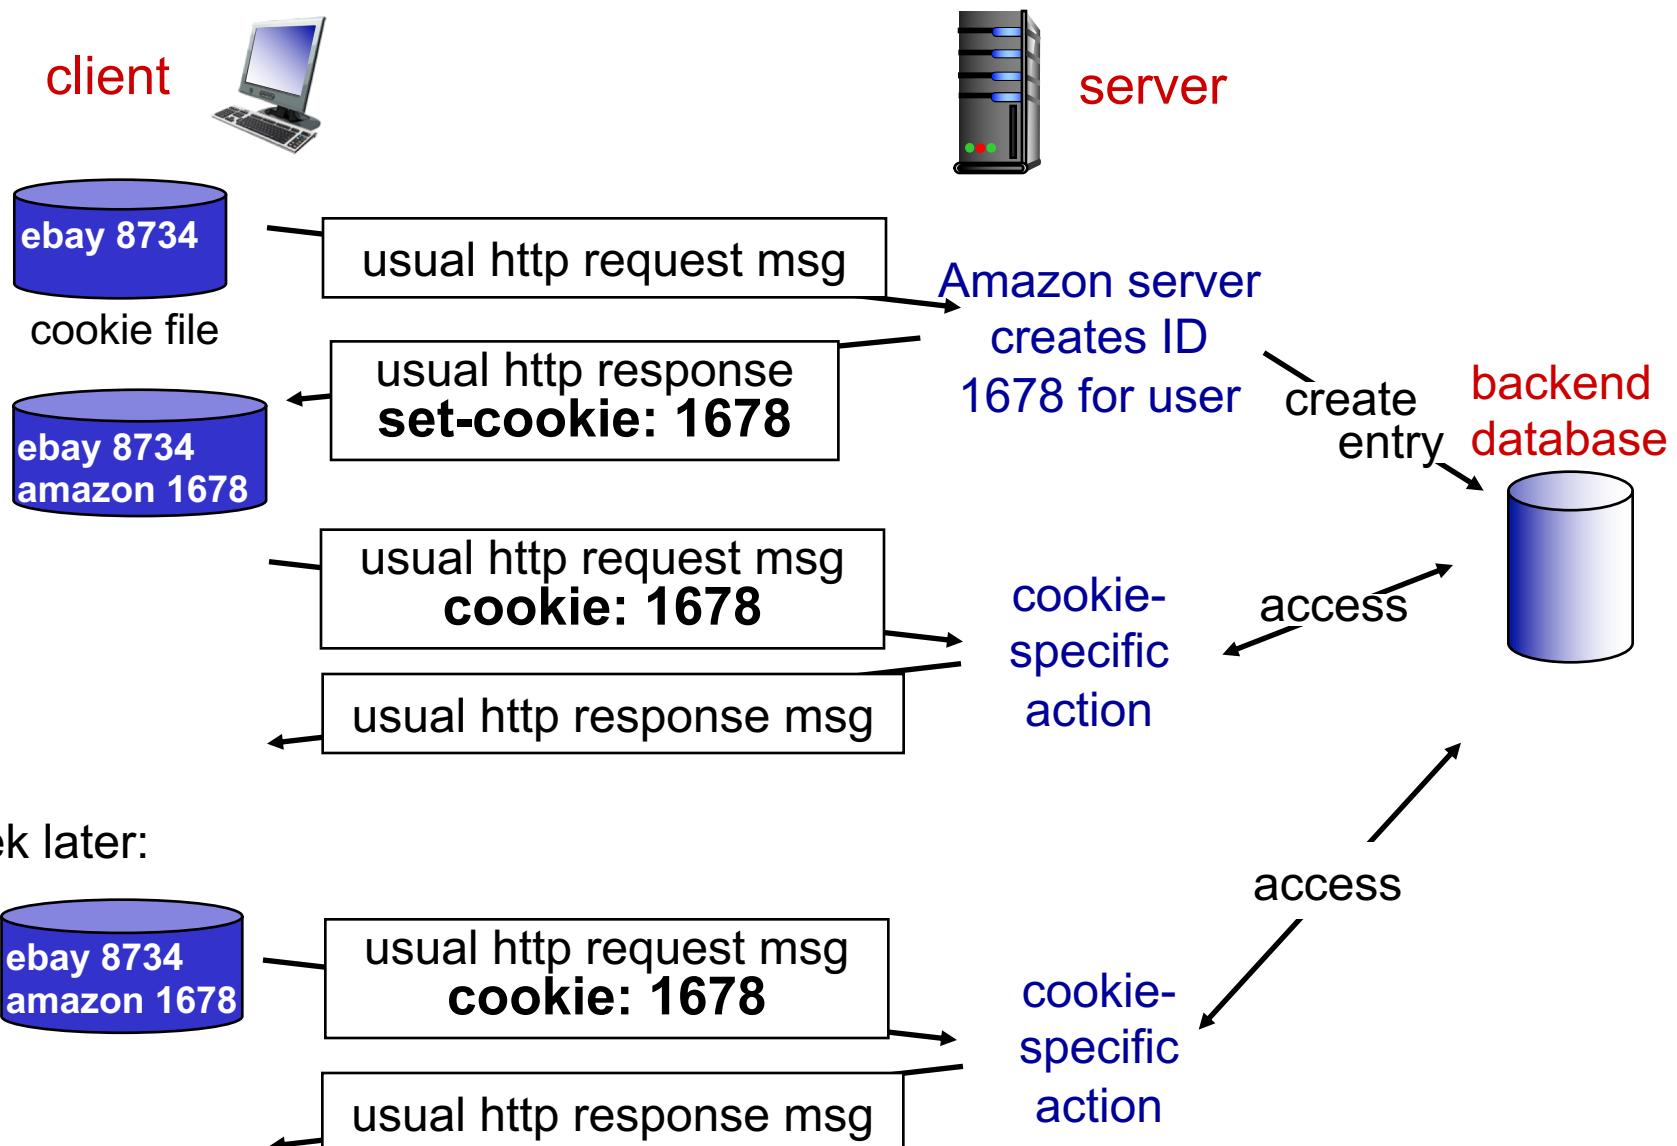
*four components:*

- 1) cookie header line of  
HTTP *response*  
message (set-cookie)
- 2) cookie header line in  
next HTTP *request*  
message (cookie)
- 3) cookie file kept on  
user's host, managed  
by user's browser
- 4) back-end database at  
Web site

**example:**

- Susan always access Internet  
from PC
- visits specific e-commerce  
site for first time
- when initial HTTP requests  
arrives at site, site creates:
  - unique ID
  - entry in backend  
database for ID

# Cookies: keeping “state” (cont.)



# Cookies (continued)

*what cookies can be used  
for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

aside  
*cookies and privacy:*

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

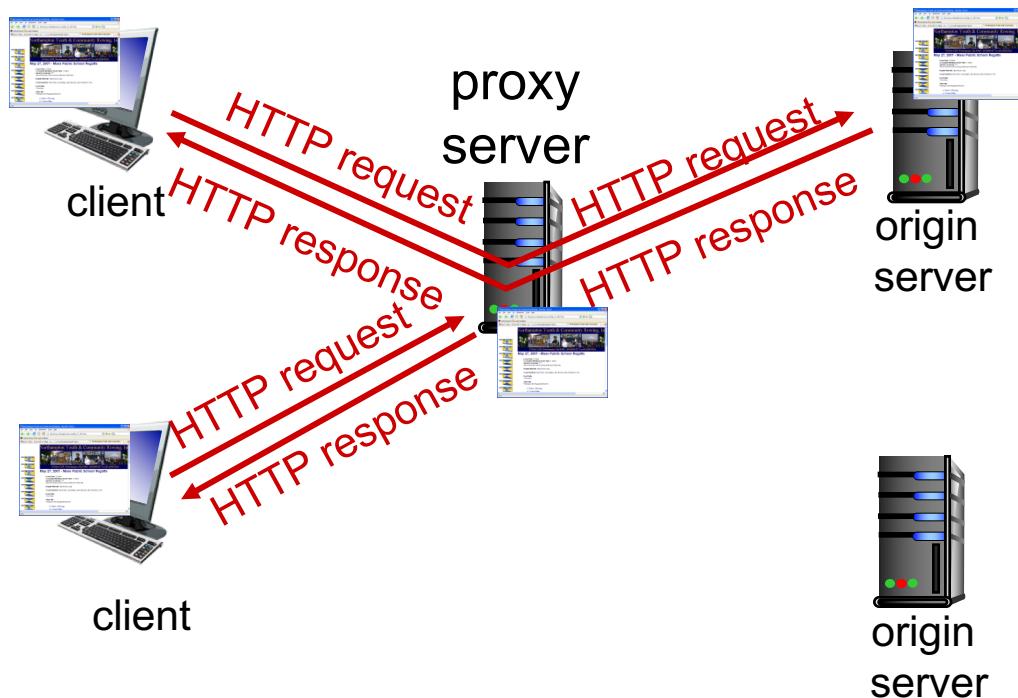
*how to keep “state”:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies create a “user session layer” on top of stateless HTTP
- cookies may be considered an invasion of privacy!

# Web caches (proxy server)

**goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache (proxy)
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client



# More about Web caching

- cache acts as both client and server
  - server for original requesting client
  - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)
- Example: Squid Web Cache

<http://www.squid-cache.org>



## *why Web caching?*

- reduce response time for client request
- reduce traffic on an institution's access link
- Increasingly important in the Internet because of CDNs (Content Distribution Networks, e.g. Google, Akamai, etc).

# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version

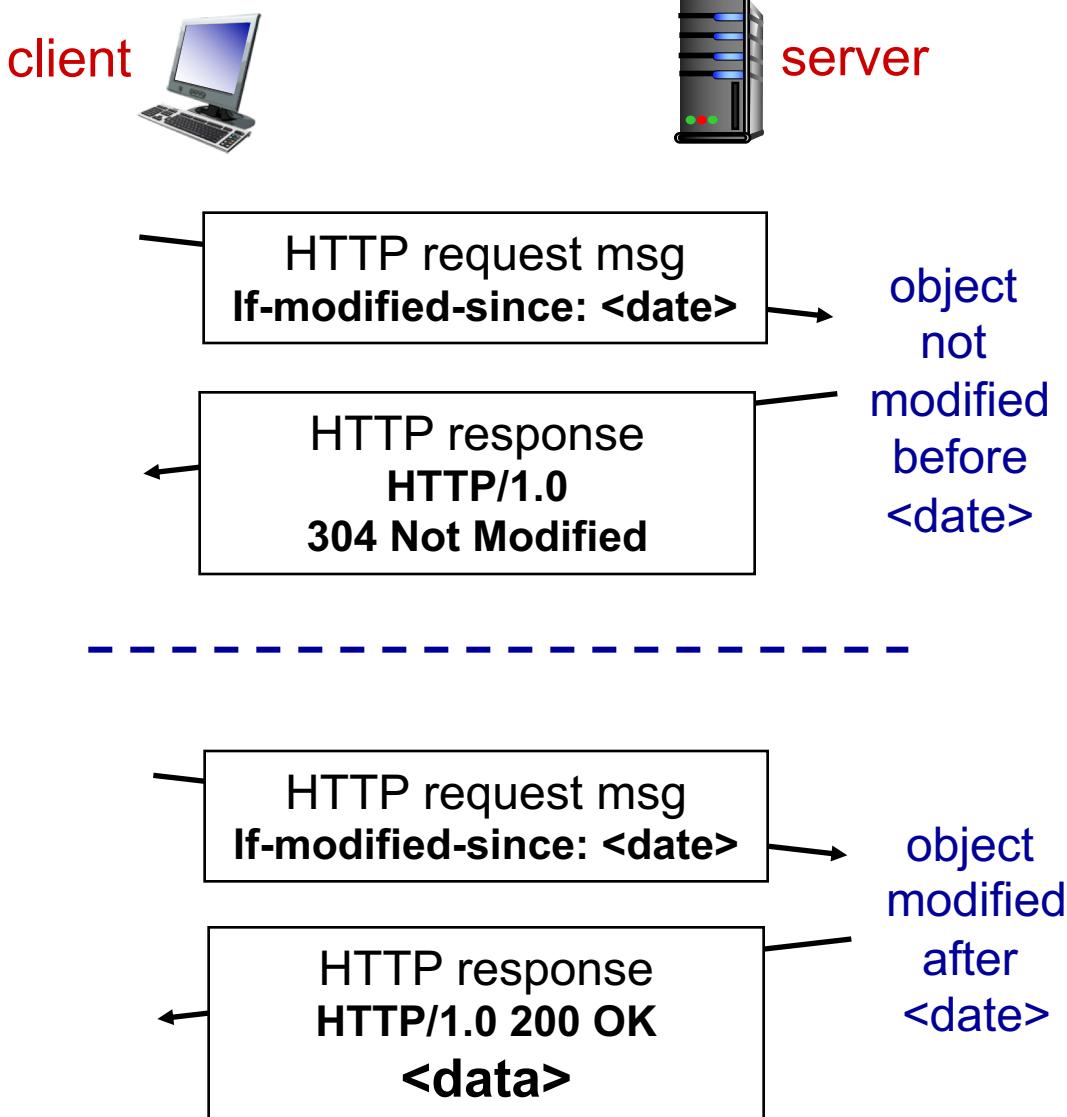
- no object transmission delay
  - lower link utilization

- **cache:** specify date of cached copy in HTTP request

**If-modified-since:**  
**<date>**

- **server:** response contains no object if cached copy is up-to-date:

**HTTP/1.0 304 Not Modified**



# T03

## Application Layer

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

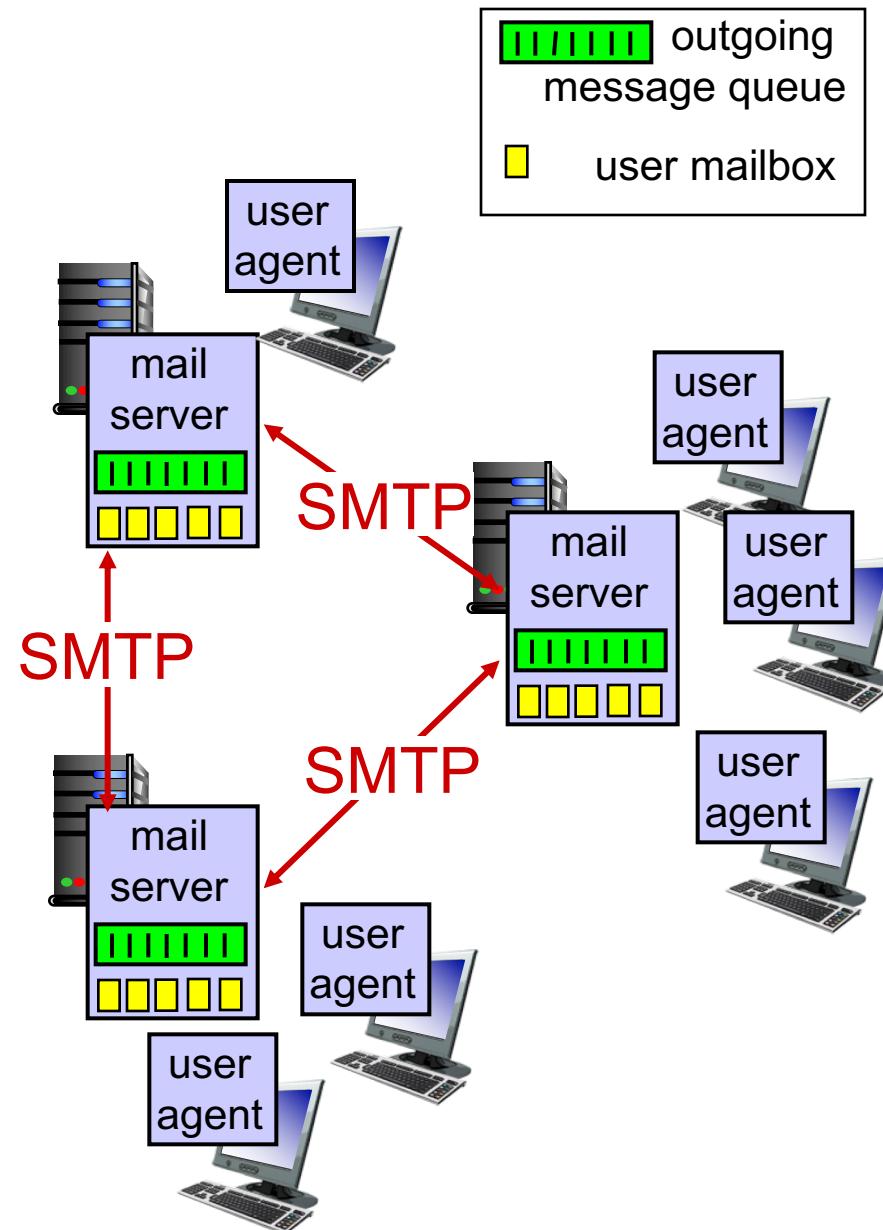
# Electronic mail

*Three major components:*

- user agents
- mail servers
- simple mail transfer protocol: SMTP

## User Agent

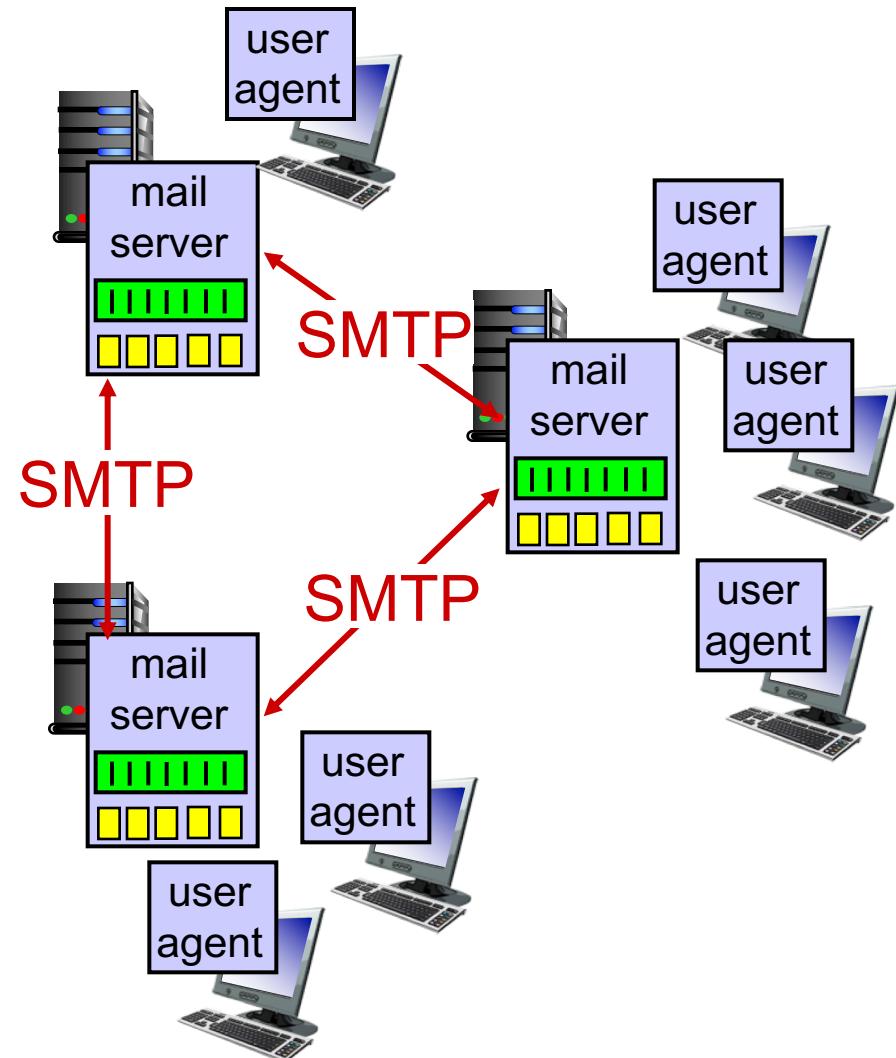
- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, Thunderbird, iPhone mail client
- outgoing, incoming messages stored on server



# Electronic mail: mail servers

## mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages
- *SMTP protocol* between mail servers to send email messages
  - client: sending mail server
  - “server”: receiving mail server
  - Every mail server is an SMTP client and server

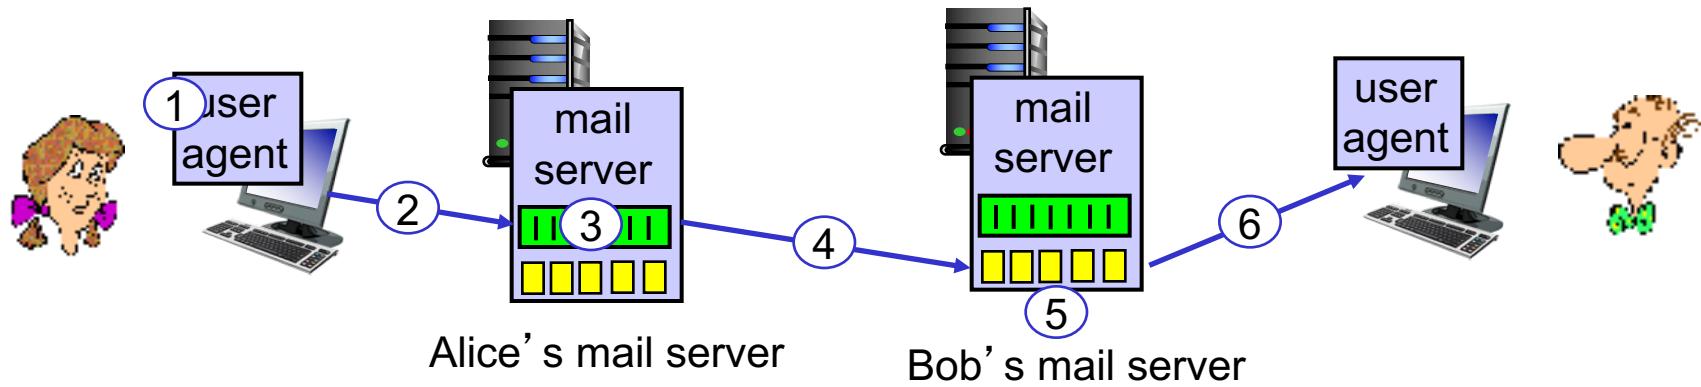


# Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- direct transfer: sending server to receiving server
- three phases of transfer
  - handshaking (greeting)
  - transfer of messages
  - closure
- command/response interaction (like HTTP)
  - commands: ASCII text
  - response: status code and phrase
- messages must be in 7-bit ASCII

# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



# Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

# Try SMTP interaction for yourself:

- `telnet smtp.dei.uc.pt 25`
- see 220 reply from server
- enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

```
[jgranjal@eden ~ $ telnet smtp.dei.uc.pt 25
Trying 193.137.203.253...
Connected to smtp.dei.uc.pt.
Escape character is '^}'.
220 smtp.dei.uc.pt ESMTP Sendmail 8.15.2/8.15.2; Sun, 29 Sep 2019 12:19:20 +0100
```

(enter other SMTP commands here – see previous example!)

# SMTP: final words

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF .CRLF to determine end of message



## *comparison with HTTP:*

- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message (e.g. body of message, images, videos, etc.)

# Mail message format

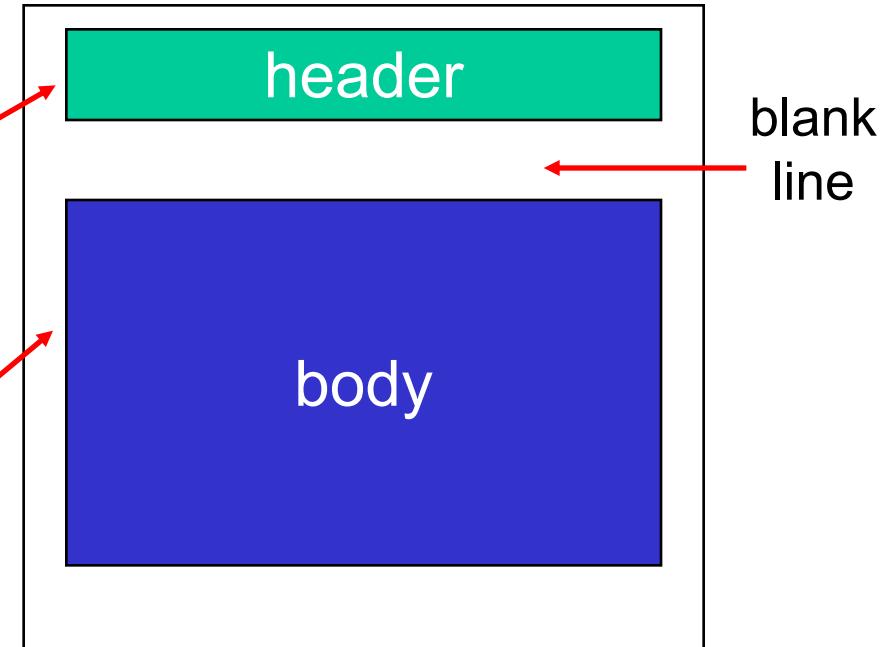
SMTP: protocol for  
exchanging email messages

RFC 822: standard for text  
message format:

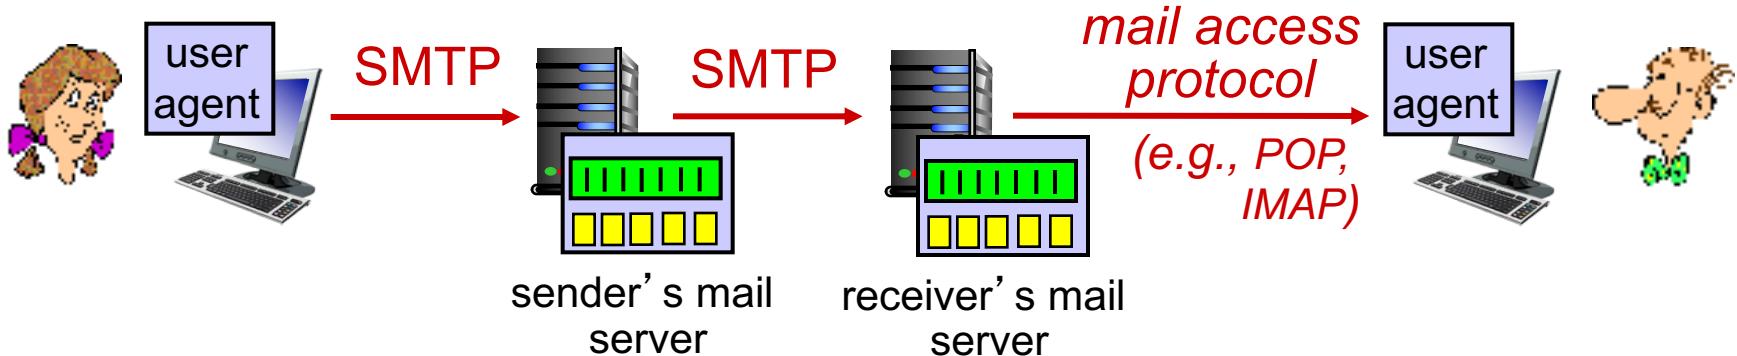
- header lines, e.g.,
  - To:
  - From:
  - Subject:

*different from SMTP MAIL  
FROM, RCPT TO:  
commands! (so called  
“envelope” headers)*

- Body: the “message”
  - ASCII characters only



# Mail access protocols



- **SMTP:** delivery/storage to receiver's server
- mail access protocol: retrieval from server
  - **POP:** Post Office Protocol [RFC 1939]: authorization, download
  - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server
  - **HTTP (web mail):** GMail, Hotmail, Yahoo! Mail, etc.

# POP3 protocol

## *authorization phase*

- client commands:
  - **user**: declare username
  - **pass**: password
- server responses
  - +OK
  - -ERR

## *transaction phase, client:*

- **list**: list message numbers
- **retr**: retrieve message by number
- **dele**: delete
- **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

# POP3 (more) and IMAP

## *more about POP3*

- previous example uses POP3 “download and delete” mode
  - Bob cannot re-read e-mail if he changes client
- POP3 “download-and-keep”: copies of messages on different clients
- POP3 is *stateless* across sessions

## *IMAP*

- keeps all messages in one place: at server
- allows user to organize messages in folders
- keeps *user state* across sessions:
  - names of folders and mappings between message IDs and folder name

# T03

## Application Layer

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

# DNS: domain name system

*people*: many identifiers:

- SSN, name, passport #

*Internet hosts, routers*:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

*Q*: how to map between IP address and name, and vice versa ?

*Domain Name System*:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
- core Internet function, implemented as application-layer protocol



# DNS: services, structure

## *DNS services*

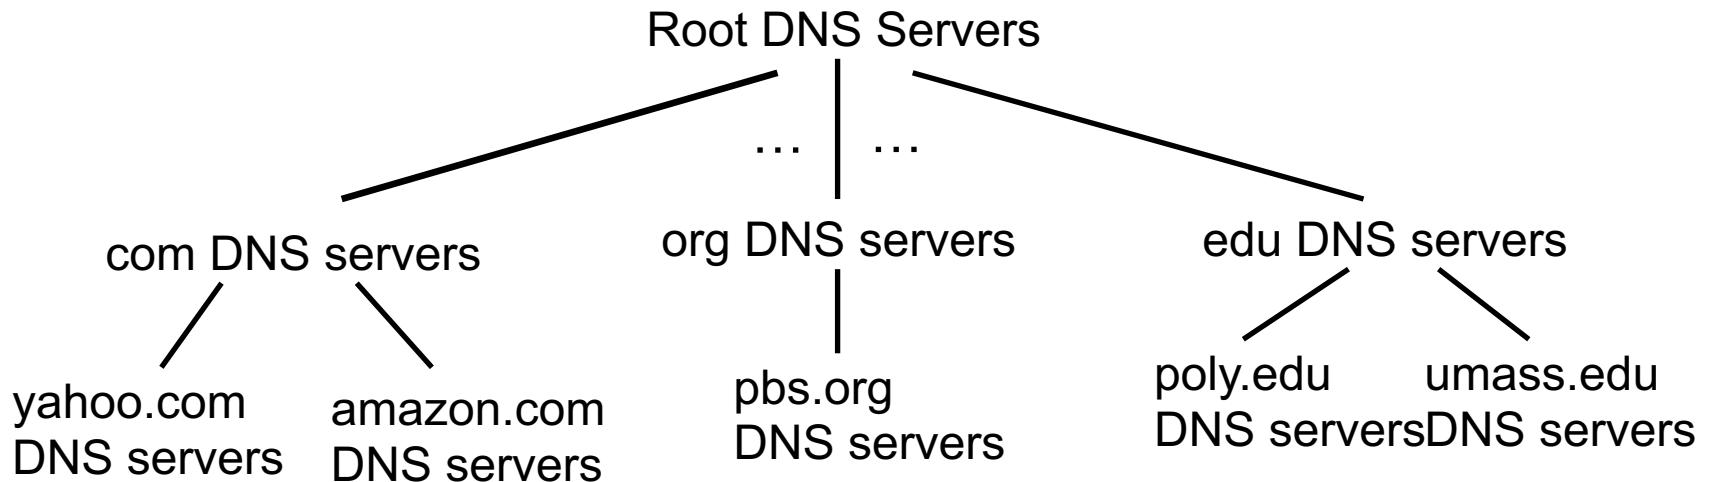
- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - Example: replicated Web servers (many IP addresses correspond to one name)

## *why not centralize DNS?*

- single point of failure
- traffic volume
- distant centralized database
- maintenance (updates)

A: *doesn't scale!*

# DNS: a distributed, hierarchical database

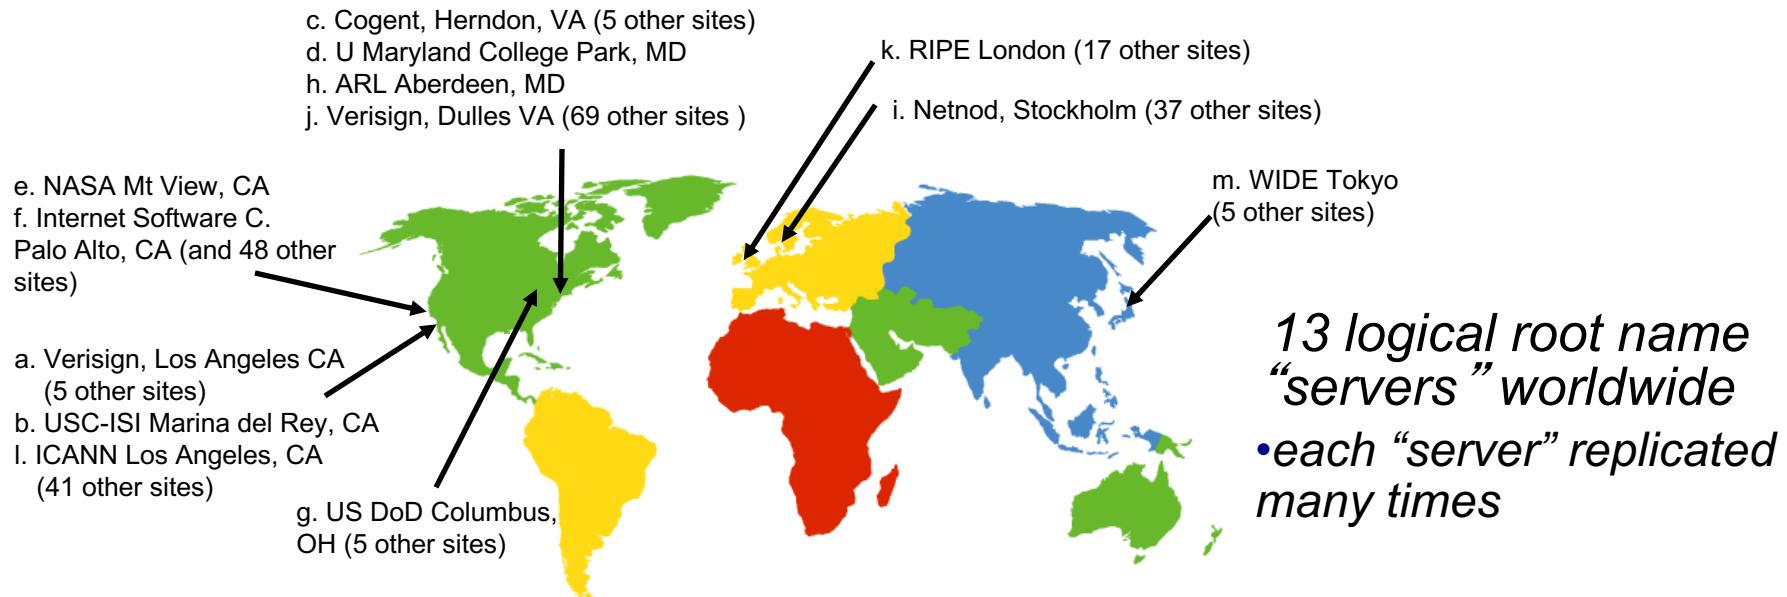


*client wants IP for www.amazon.com; 1<sup>st</sup> approximation:*

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

# DNS: root name servers

- contacted by local name server that can not resolve name
- root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server



- See also: [Root Zone Database \(IANA\)](#)

# TLD, authoritative servers

## *top-level domain (TLD) servers:*

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

## *authoritative DNS servers:*

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Local DNS name server

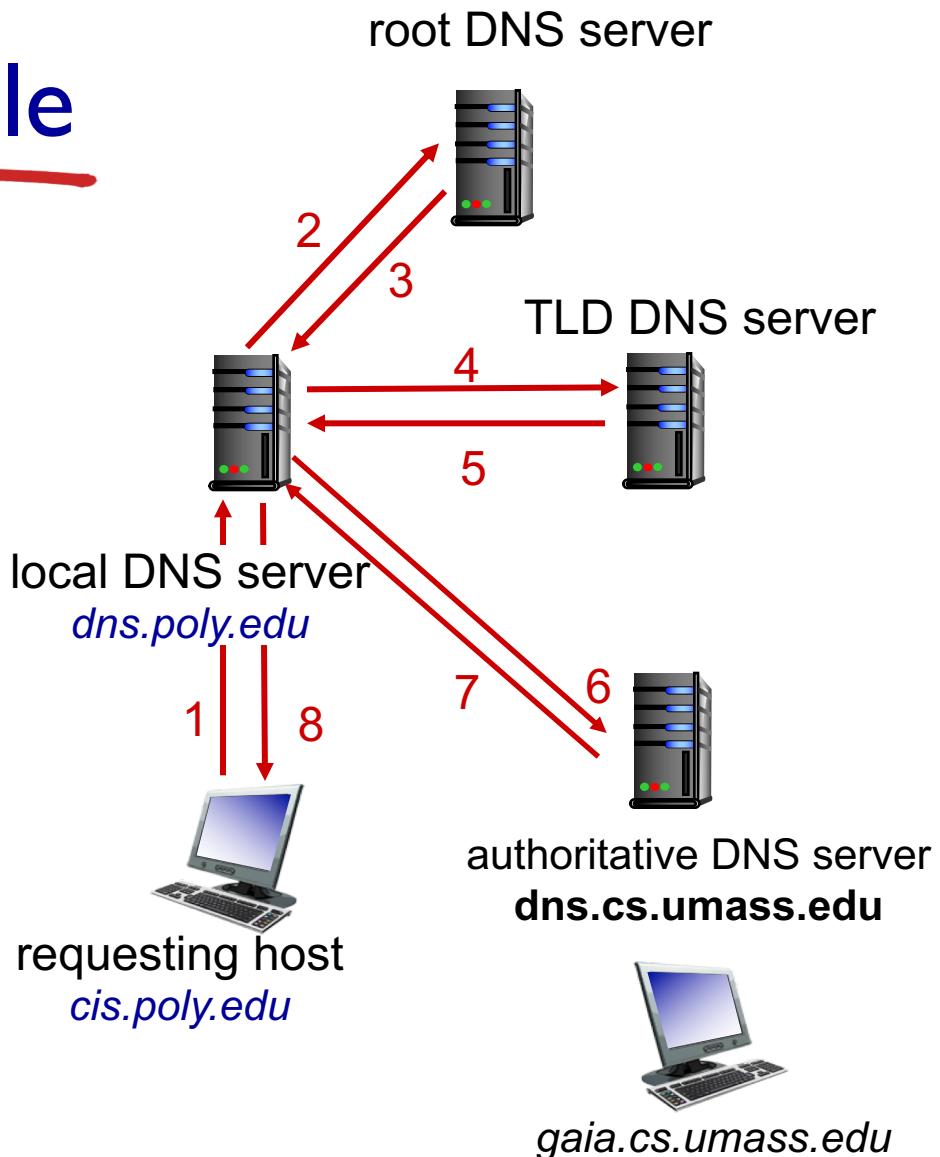
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
  - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date!)
  - acts as proxy, forwards query into hierarchy

# DNS name resolution example

- host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

## *iterated query:*

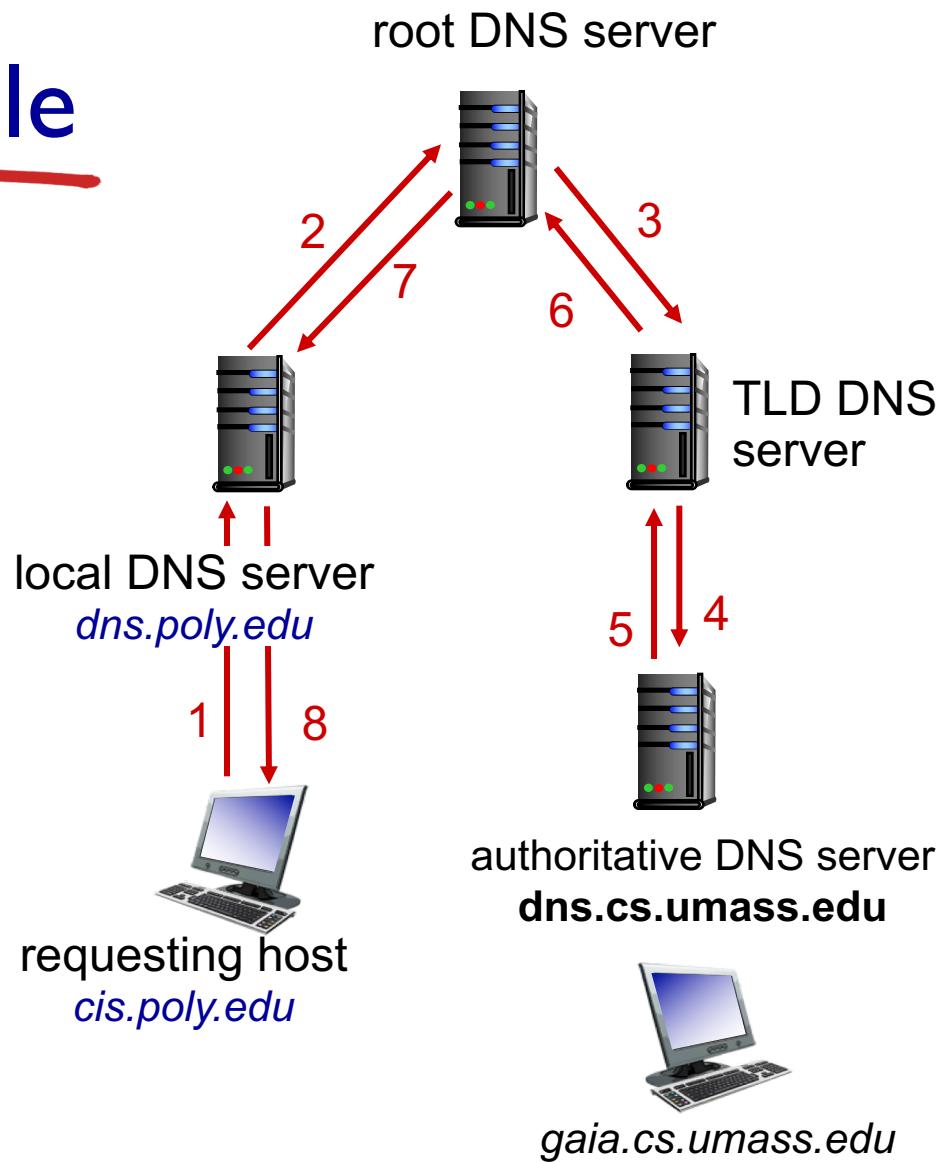
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

*recursive query:*

- puts burden of name resolution on contacted name server



# DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
  - cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - thus root name servers not often visited
- cached entries may be *out-of-date* (best effort name-to-address translation!)
  - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
  - RFC 2136

# DNS records

**DNS:** distributed database storing resource records (**RR**)

RR format: `(name, value, type, ttl)`

## type=A

- **name** is hostname
- **value** is IP address

## type=NS

- **name** is domain (e.g.,  
foo.com)
- **value** is hostname of  
authoritative name  
server for this domain

## type=CNAME

- **name** is alias name for some  
“canonical” (the real) name
- `www.ibm.com` is really  
`servereast.backup2.ibm.com`
- **value** is canonical name

## type=MX

- **value** is name of mailserver  
associated with **name**

# Inserting records into DNS

- example: new startup “Network Utopia”
- register name `networkutopia.com` at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative name server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:  
`(networkutopia.com, dns1.networkutopia.com, NS)`  
`(dns1.networkutopia.com, 212.212.212.1, A)`
- create authoritative server type A record for  
`www.networkutopia.com`; type MX record for  
`networkutopia.com`

# T03

## Application Layer

2.1 principles of network  
applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

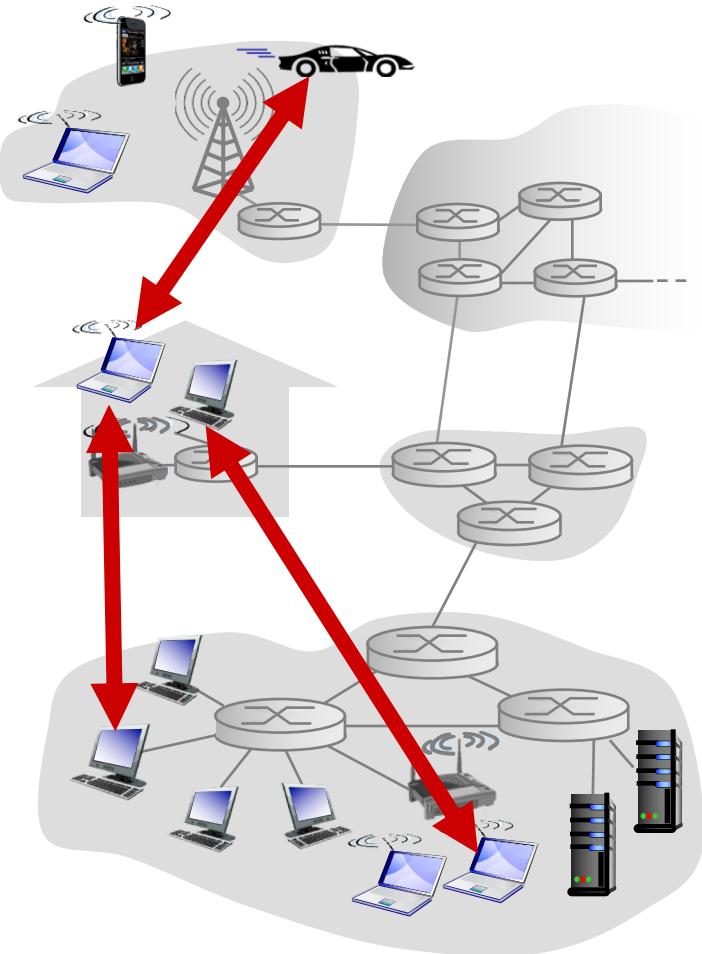
2.6 video streaming and  
content distribution  
networks

# Pure P2P architecture

- no need to rely on always-on server
- arbitrary end systems communicate directly
- peers are intermittently connected and change IP addresses

*examples:*

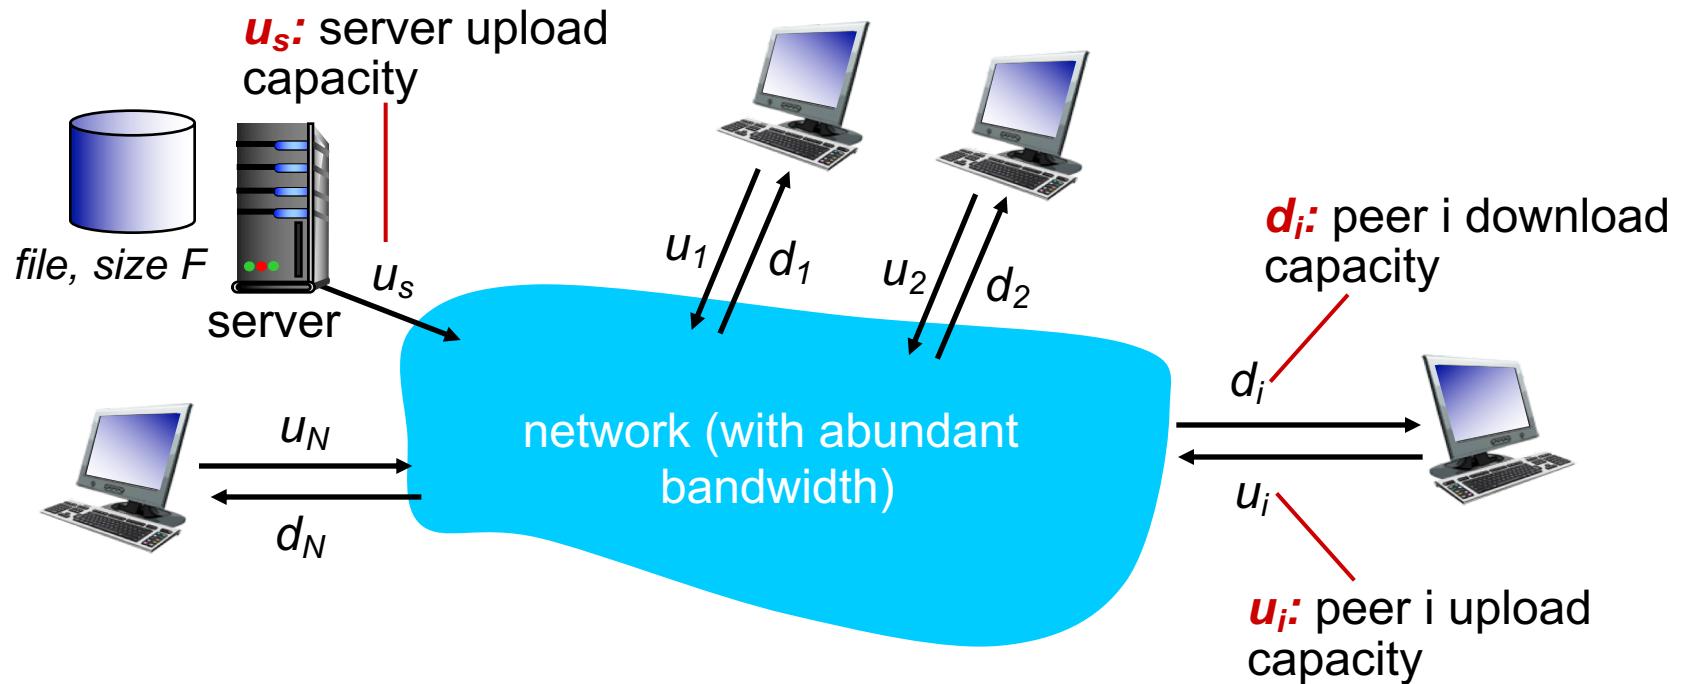
- file distribution (BitTorrent)
- VoIP (Skype)



# File distribution: client-server vs P2P

Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?

- peer upload/download capacity is limited resource



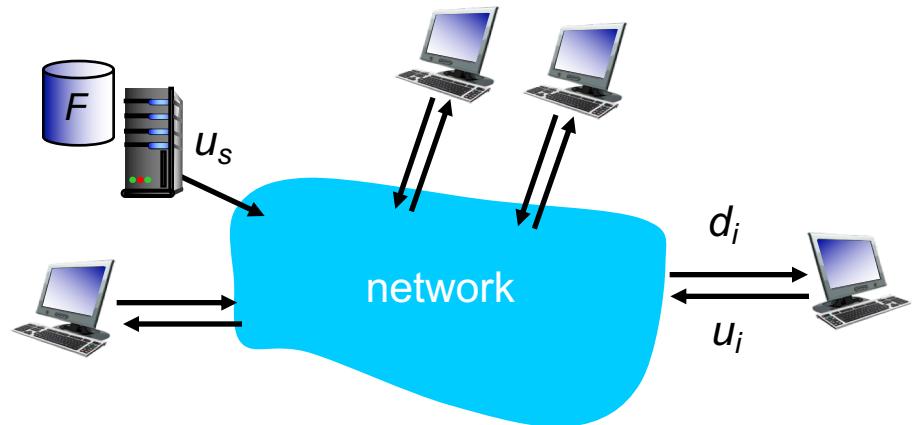
# File distribution time: client-server

- **server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$

- **client:** each client must download file copy

- $d_{min}$  = min client download rate
- min client download time:  $F/d_{min}$



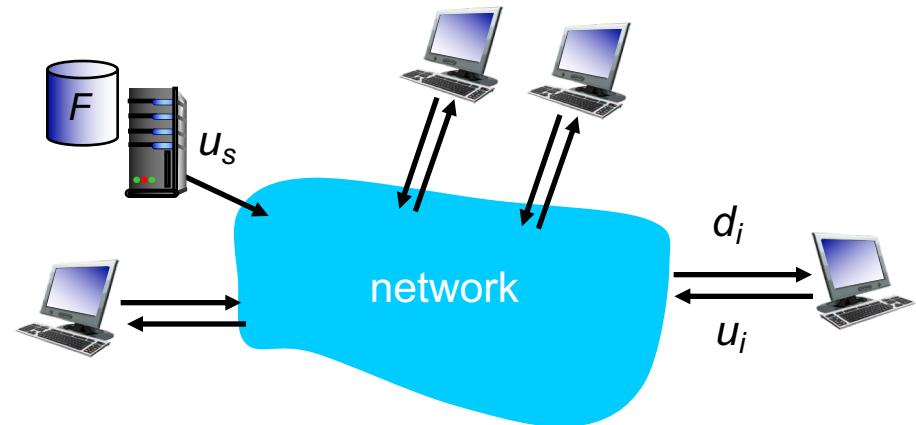
*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$  (number of peers)

# File distribution time: P2P

- **server transmission:** must upload at least one copy (at the beginning only the server has the file)
  - time to send one copy:  $F/u_s$



- **client:** each client must download file copy
  - min client download time:  $F/d_{\min}$
- **clients:** as aggregate must download  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$

time to distribute  $F$   
to  $N$  clients using  
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

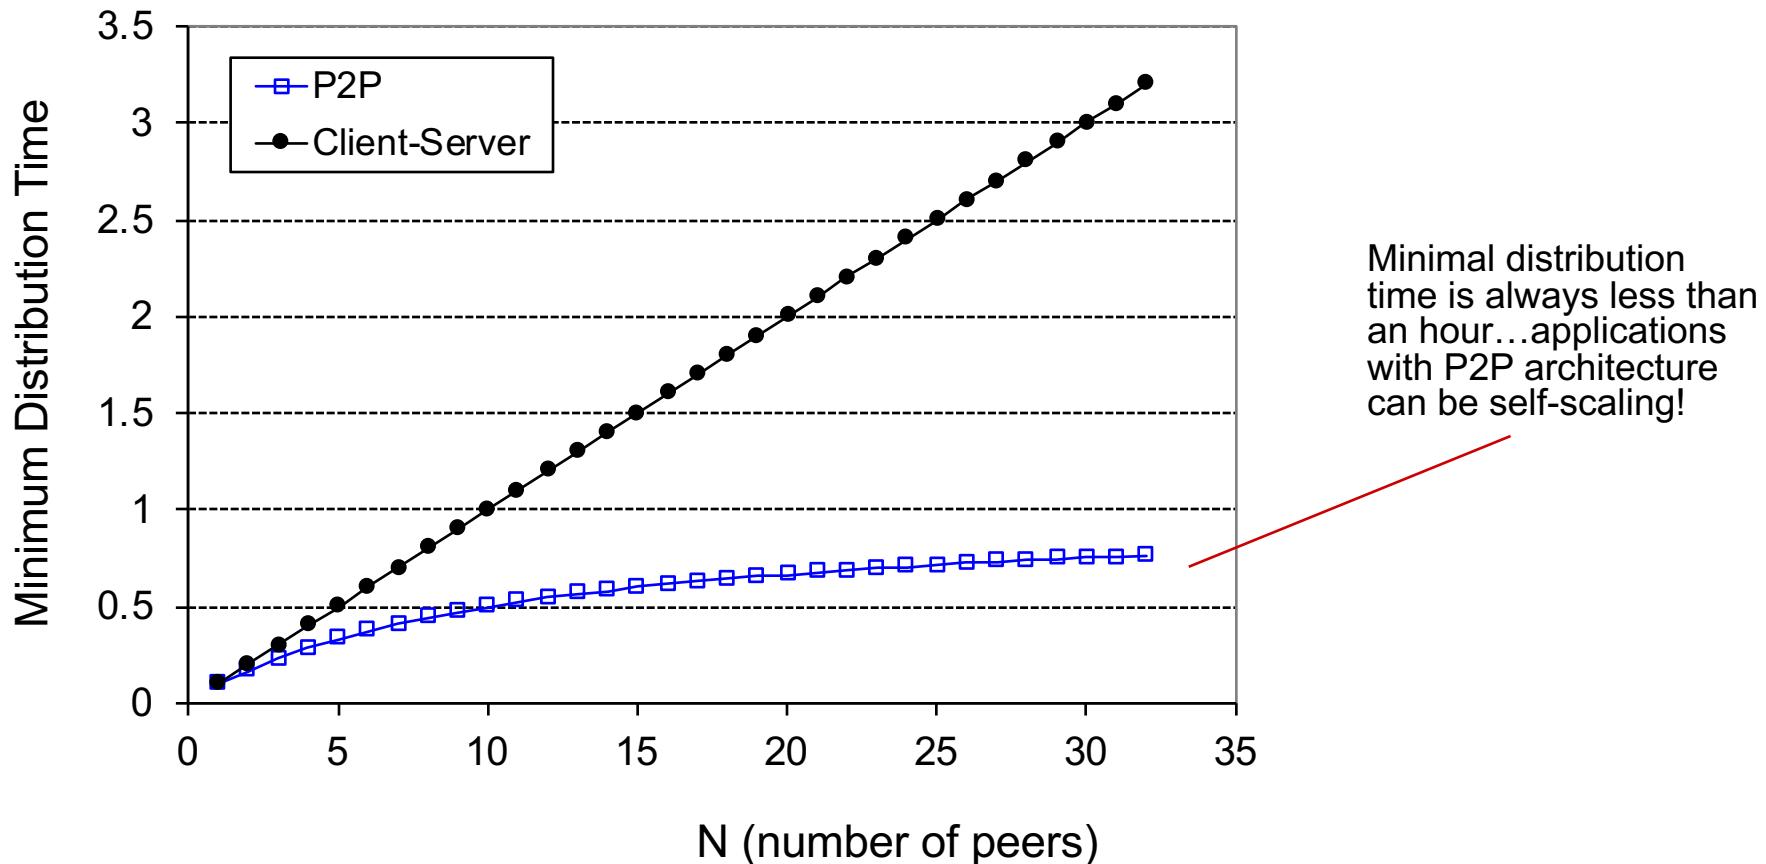
increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

(lower bound is achievable with a scheme where each peer can redistribute a bit as soon as it receives it, and a good approximation when redistributing chunks of the file)

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour (a peer can transmit the file in 1 hour),  
 $u_s = 10u$  (server transmission rate is 10 x the peer upload rate),  $d_{min} \geq u_s$

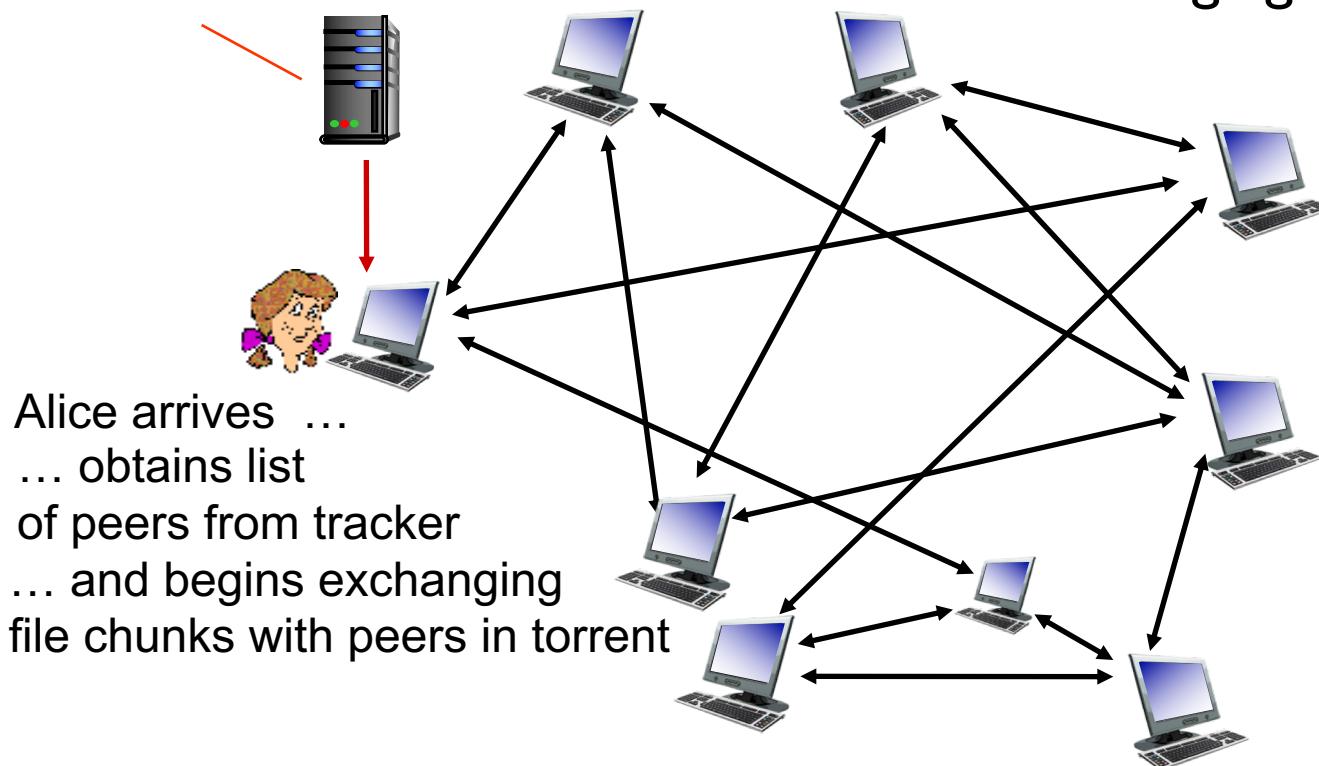


# P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

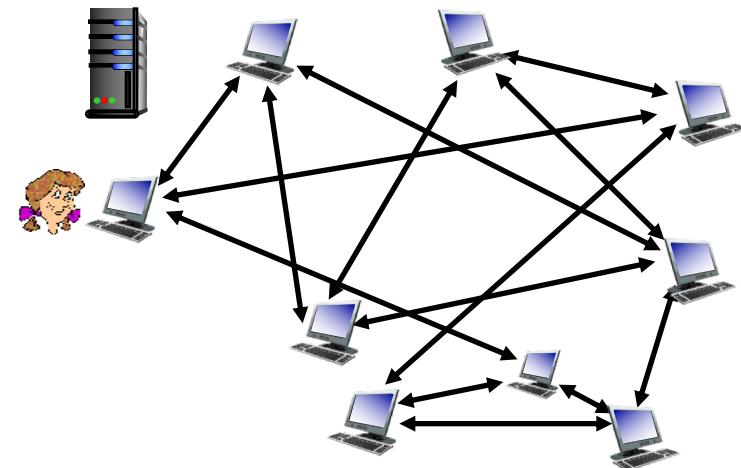
*tracker*: tracks peers  
participating in torrent

*torrent*: group of peers  
exchanging chunks of a file



# P2P file distribution: BitTorrent

- peer joining torrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- *churn*: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



# BitTorrent: requesting, sending file chunks

## *requesting chunks:*

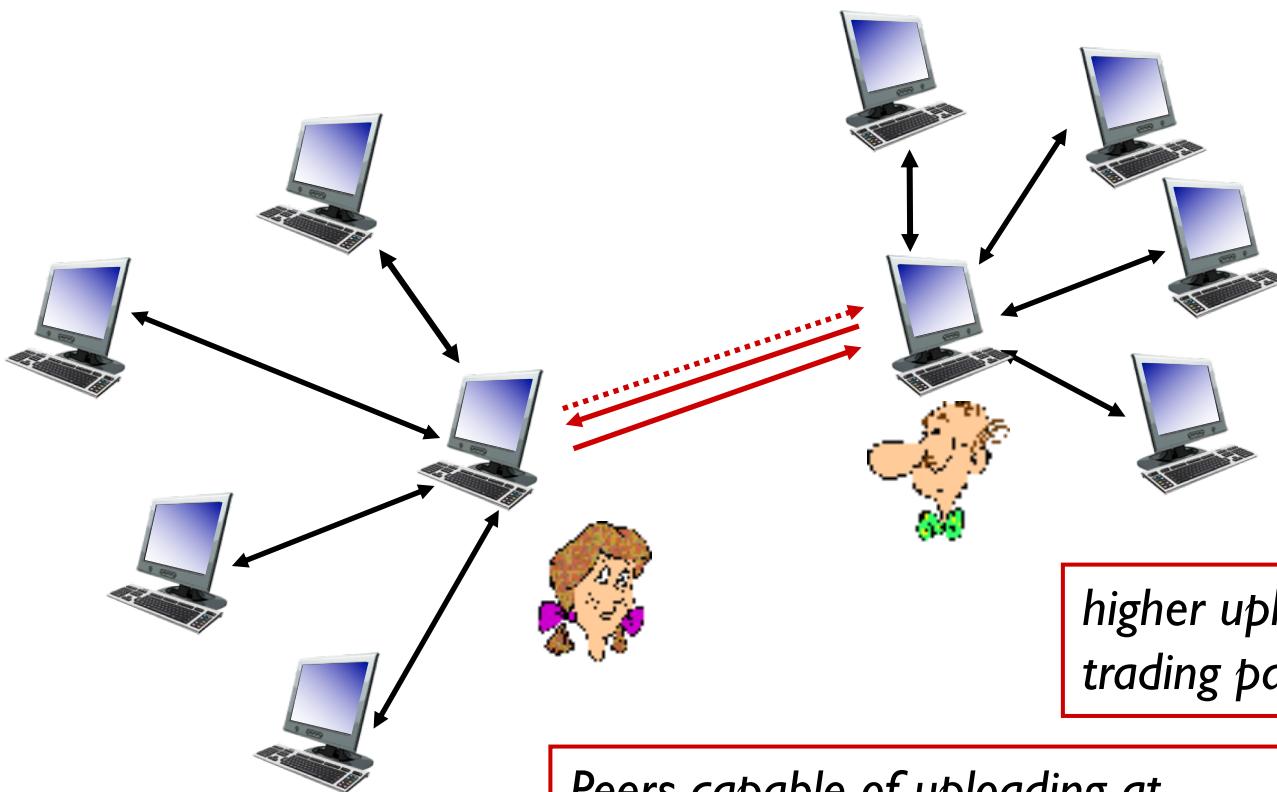
- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

## *sending chunks: tit-for-tat*

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: also randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



*“tit-for-tat” incentive mechanism for trading promotes cooperation and fairness, is at the heart of BitTorrent!*

*higher upload rate: find better trading partners, get file faster !*

*Peers capable of uploading at compatible rates tend to find each other*

# T03

## Application Layer

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

# Video Streaming and CDNs: context

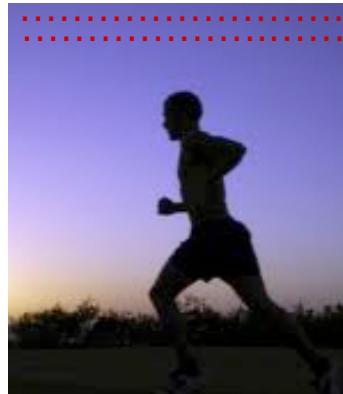
- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution:* distributed, application-level infrastructure



# Multimedia: video

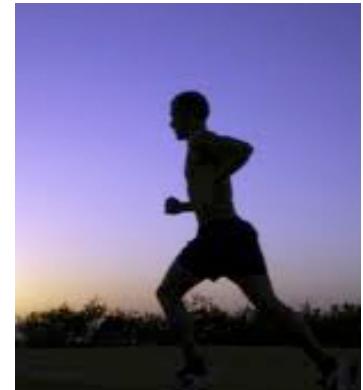
- video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- digital image: array of pixels
  - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$

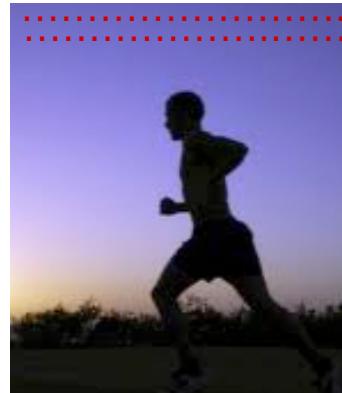


frame  $i+1$

# Multimedia: video

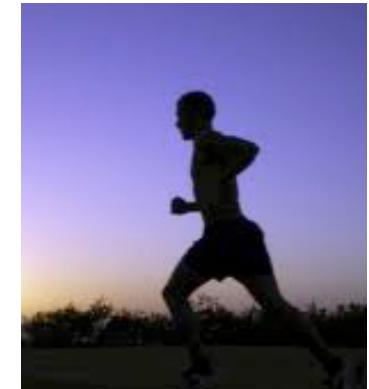
- **CBR: (constant bit rate):**  
video encoding rate fixed
- **VBR: (variable bit rate):**  
video encoding rate changes  
as amount of spatial,  
temporal coding changes
- **examples:**
  - MPEG I (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



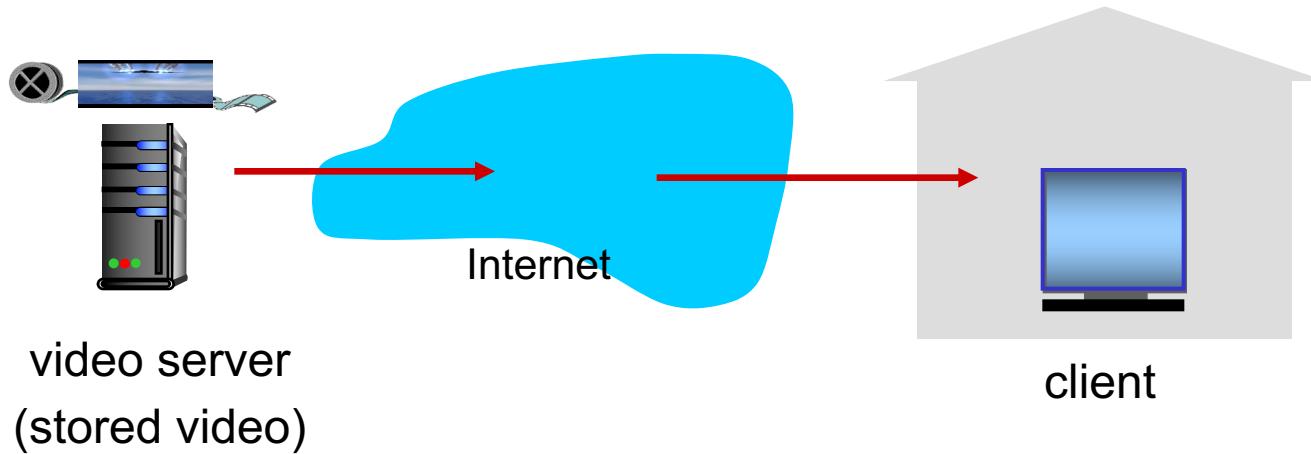
frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



# Streaming stored video:

simple scenario:



# Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **server:**
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - *manifest file*: provides URLs for different chunks
- **client:**
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “*intelligence*” at client: client determines
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)
- Examples: MPEG-DASH in YouTube and Netflix



# Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users around the world?
- Example: a billion hours of video watched on [YouTube](#) every day!
- **Option 1:** single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients (bottleneck links)
  - multiple copies of video (e.g. popular videos) sent over outgoing link, which Internet video company has to pay

....quite simply: this solution **doesn't scale!**

# Content distribution networks

- ***challenge:*** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users around the world?
- ***option 2:*** store/serve multiple copies of videos at multiple geographically distributed sites (***CDN***). Two different server placement strategies:
  - ***enter deep:*** push CDN servers deep into many access networks
    - get close to end users
    - used by Akamai, 1700 locations
  - ***bring home:*** smaller number (10's) of larger clusters in POPs near (but not within) access networks (Tier-1 ISPs)
    - used by Limelight

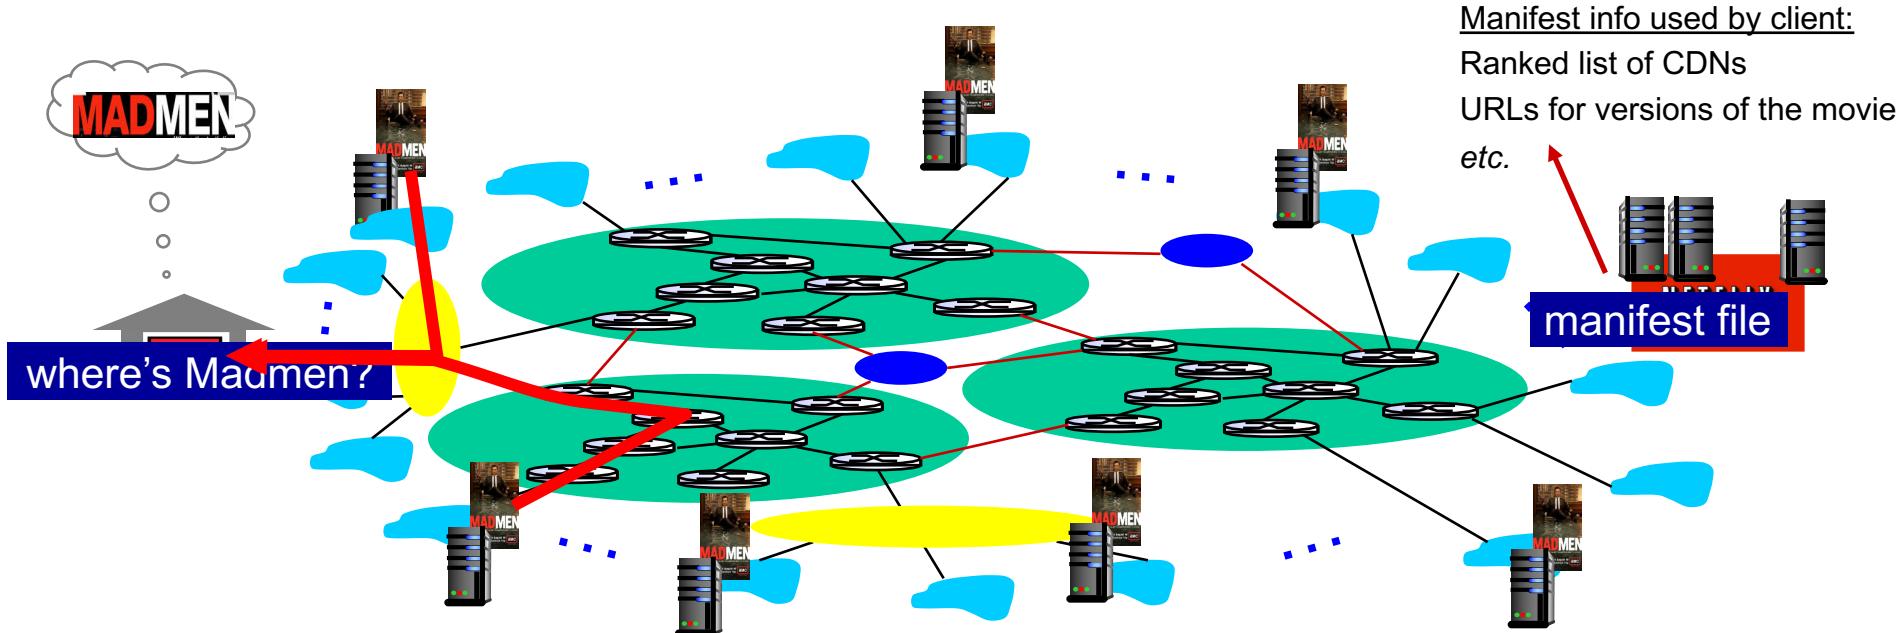
# Content distribution networks

- Content Distribution (or Delivery) Networks (DNS) may be *private* or *third-party*
- A *private* CDN is owned by the content provider, e.g. Google's CDN distributes YouTube Videos
- A *third-party* CDN distributed content on behalf of multiple content-providers, e.g. Limelight or Akamai for Netflix and Hulu



# Content Distribution Networks (CDNs)

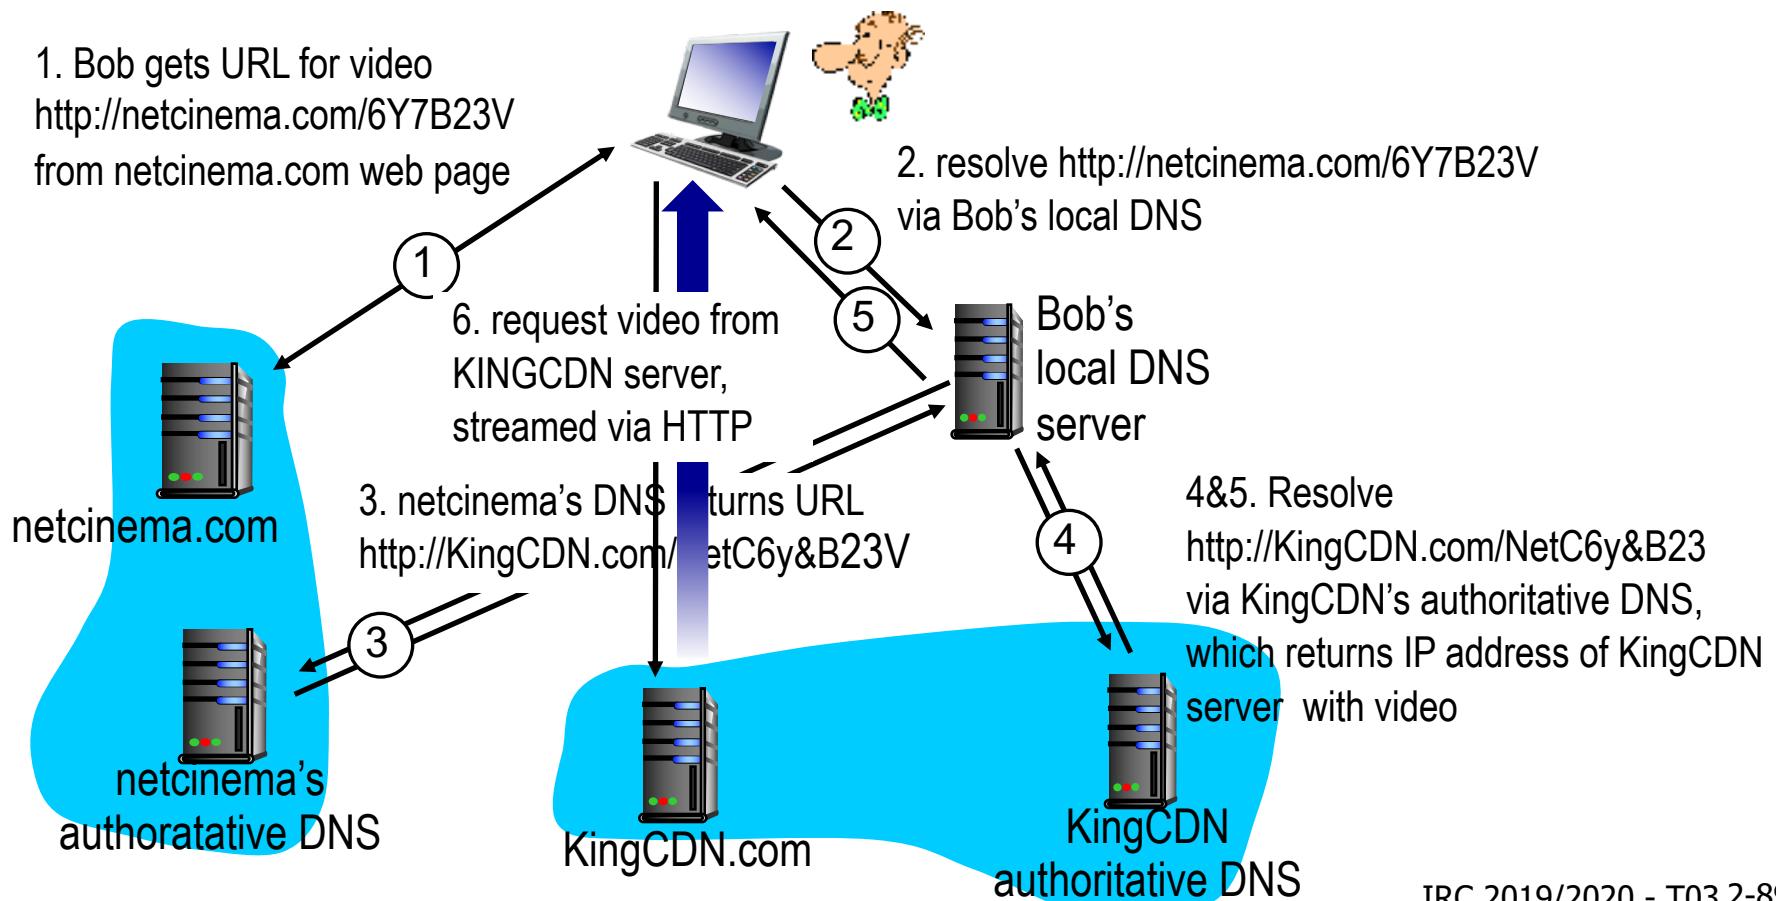
- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



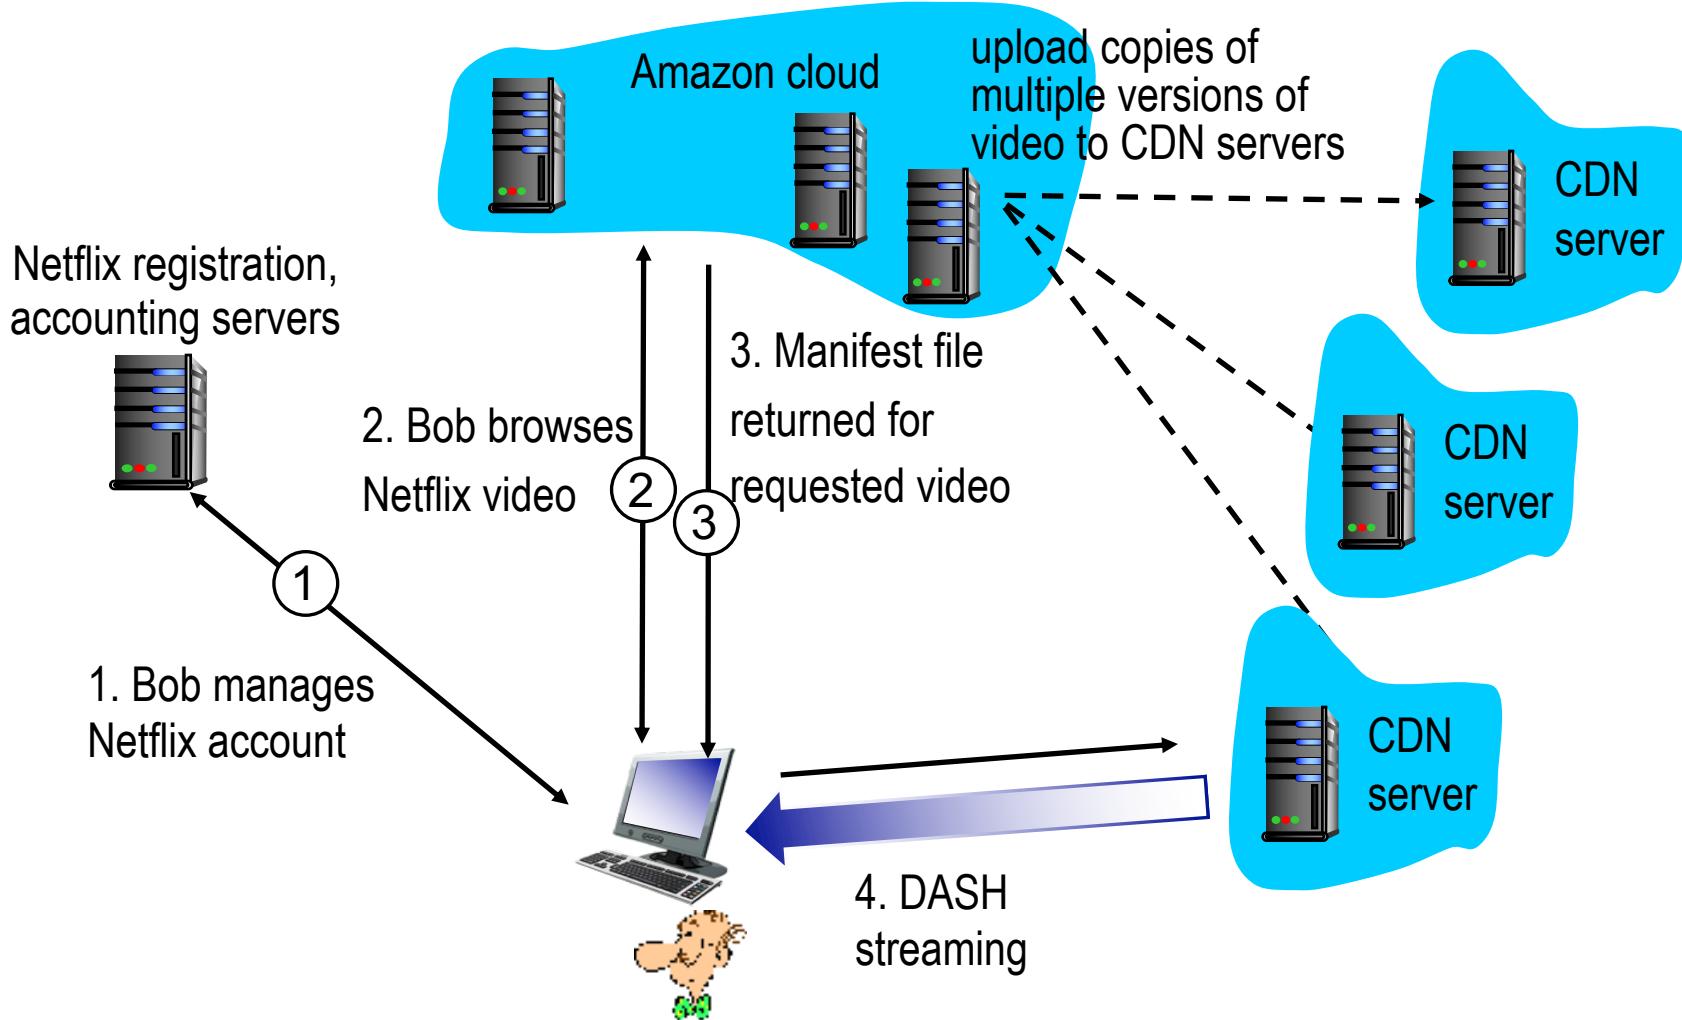
# CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



# Case study: Netflix



# T03 Application Layer

*our study of network apps now complete!*

- application architectures
  - client-server
  - P2P
- application service requirements:
  - reliability, bandwidth, delay
- Internet transport service model
  - connection-oriented, reliable: TCP
  - unreliable, datagrams: UDP
- specific protocols:
  - HTTP
  - SMTP, POP, IMAP
  - DNS
  - P2P: BitTorrent
- video streaming, CDNs

# T03 Application Layer

*most importantly: learned about protocols!*

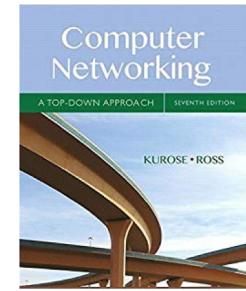
- typical request/reply message exchange:
  - client requests info or service
  - server responds with data, status code
- message formats:
  - *headers*: fields giving info about data
  - *data*: info(payload) being communicated

*important themes:*

- control vs. messages
  - in-band, out-of-band
- centralized vs. decentralized
- stateless vs. stateful
- reliable vs. unreliable message transfer
- “complexity at network edge”

# T03: Bibliography

J. Kurose and K. Ross, "Computer Networking - a top-down approach", Pearson. Chapter 2: Application Layer



# Introdução às Redes e Comunicações

## 2019/2020

T03

Application Layer  
Extra material

Jorge Granjal  
University of Coimbra



# T03: Netflix Open Connect

- Open Connect is a Netflix open source project, over time the company is moving traffic from third-party CDN providers
- ISP partner with Netflix to localize traffic and minimize data usage of transit provider

Quick Guide: What Is Netflix Open Connect

Netflix: Welcome to Open Connect



# T03: Tor and Onion Routing

- The Tor network supports anonymization of web browsing
- Tor uses Onion Routing, where messages are encapsulated in layers of encryption

## Onion Routing

The dark side of the web -- exploring darknets

